PROJECT REPORT

On

Examcraft.AI

Submitted by

Meetkumar Mistry (IU2241050071)
Smit Patel (IU2241050096)
Rushil Patel (IU2241050123)

In fulfillment for the award of the degree

Of

BACHELOR OF TECHNOLOGY

In

COMPUTER ENGINEERING

INSTITUTE OF TECHNOLOGY AND ENGINEERING

INDUS UNIVERSITY CAMPUS, RANCHARDA, VIA-THALTEJ

AHMEDABAD-382115, GUJARAT, INDIA,

WEB: www.indusuni.ac.in

MAY 2026

PROJECT REPORT

ON

Examcraft.AI

AT

In the partial fulfillment of the requirement
for the degree of
Bachelor of Technology
in
Computer Engineering

PREPARED BY

Meetkumar Mistry (IU2241050071)
Smit Patel (IU2241050096)
Rushil Patel(IU2241050123)

UNDER GUIDANCE OF

Internal Guide
Shreya Dubey
Lecturer,
Department of Computer Engineering,
I.I.T.E, Indus University ,Ahmedabad

SUBMITTED TO

INSTITUTE OF TECHNOLOGY AND ENGINEERING
INDUS UNIVERSITY CAMPUS, RANCHARDA, VIA-THALTEJ
AHMEDABAD-382115, GUJARAT, INDIA,
WEB: www.indusuni.ac.in
MAY 2026

CANDIDATE'S DECLARATION

I declare that final semester report entitled "**EXAMCRAFT.AI**" is my own work conducted under the supervision of the guide **SHREYA DUBEY**.

I further declare that to the best of my knowledge, the report for B.Tech final semester does not contain part of the work which has been submitted for the award of B.Tech Degree either in this university or any other university without proper citation.

_____

Candidate's Signature

Meetkumar Mistry(IU2241050071)

_____

Candidate's Signature

Smit Patel(IU2241050096)

_____

Candidate's Signature

Rushil Patel(IU2241050123)

_____

Guide : Shreya Dubey
Lecturer
Department of Computer Engineering,
Indus Institute of Technology and Engineering
INDUS UNIVERSITY– Ahmedabad,
State: Gujarat

INDUS INSTITUTE OF TECHNOLOGY AND ENGINEERING
COMPUTER ENGINEERING
2025 -2026

CERTIFICATE

Date: 24/04/2026

This is to certify that the project work entitled "**EXAMCRAFT.AI"** has been carried out by **MEETKUMAR MISTRY, SMIT PATEL and RUSHIL PATEL** under my guidance in partial fulfillment of degree of Bachelor of Technology in **COMPUTER ENGINEERING (Final Year)** of Indus University, Ahmedabad during the academic year 2025 - 2026

SHREYA DUBEY
Lecturer
Department of Computer Engineering,
I.T.E, Indus University
Ahmedabad

Dr. SEEMA MAHAJAN
Head of the Department,
Department of Computer Engineering,
I.T.E, Indus University
Ahmedabad

ACKNOWLEDGEMENT

We take immense pleasure in expressing our heartfelt gratitude to all those who have contributed to the successful completion of this project, "EXAMCRAFT.AI: An AI-Powered Academic Examination Paper Generation System."

First and foremost, We express our sincere and deep gratitude to our Faculty Guide for their invaluable guidance, consistent encouragement, and constructive feedback throughout the entire duration of this project. Their academic expertise and practical insights Ire instrumental in shaping the direction and quality of our work.

We extend our heartfelt thanks to the Head of the Department of Computer Engineering, Indus University, for providing the necessary infrastructure, academic support, and an enabling environment that made the completion of this project possible.

We are also grateful to the entire faculty of the Department of Computer Engineering for their continuous academic mentorship and the strong foundation in software engineering, data structures, artificial intelligence, and the technologies that made the technical implementation of this project achievable.

Special acknowledgement is extended to the developers and maintainers of the open-source technologies employed in this project, including Flask, React, Groq API, Supabase, ReportLab, Pydantic, and the LLaMA 3.3 language model. The availability of these robust tools significantly accelerated the development process.

We are thankful to our friends and fellow students for their motivation and peer support. Finally, we are deeply grateful to our families for their unwavering faith, patience, and encouragement, which sustained us through the challenges of this academic undertaking.

- Meetkumar Mistry(IU2241050071)
-Smit Patel(IU2241050096)
-Rushil Patel(IU2241050123)
Computer Engineering

TABLE OF CONTENT

ABSTRACT

ExamCraft.AI is a full-stack, AI-powered web application that automates the generation of structured, syllabus-aligned examination papers. It addresses the time-intensive and expertise-driven nature of traditional paper setting by intelligently generating balanced question papers based on uploaded course material and pedagogical frameworks such as Bloom's Taxonomy.

The system follows a decoupled architecture with a React (Vite, Tailwind CSS) frontend and a Flask-based RESTful backend. Users upload study material in PDF, DOCX, or TXT formats, which is processed and passed to the Groq API using the LLaMA 3.3-70B-Versatile model. The AI generates exam papers according to user-defined parameters, including section structure, marks distribution, difficulty levels, and Bloom's cognitive categories (Remembering to Creating). Additional outputs such as answer keys, marking schemes, and estimated response times are also supported.

A key contribution is the integration of Pydantic-based schema validation applied directly to AI outputs, ensuring structural correctness and eliminating issues like malformed JSON or missing fields. The system also enables selective regeneration of individual questions, improving flexibility without requiring full paper regeneration.

Data persistence is managed through Supabase, which provides PostgreSQL storage, JWT-based authentication, and cloud file handling. Generated papers can be exported as password-protected PDFs using ReportLab or as DOCX files, ensuring both security and usability. API rate limiting via Flask-Limiter enhances system reliability under load.

Evaluation through functional testing and performance analysis shows that ExamCraft.AI consistently produces well-structured papers with high adherence to specified configurations. The system demonstrates strong potential for adoption in educational institutions and digital learning platforms.

Keywords: *Examination Paper Generation, Artificial Intelligence, Bloom's Taxonomy, Large Language Models, LLaMA, Flask, React, Pydantic, Supabase, Educational Technology.*

LIST OF FIGURES

LIST OF TABLES

ABBREVIATION

Abbreviation	Full Form
AI	Artificial Intelligence
AIEd	Artificial Intelligence in Education
AQG	Automated Question Generation
API	Application Programming Interface
AES	Advanced Encryption Standard
Auth	Authentication / Authorization
B2B	Business-to-Business
B2C	Business-to-Consumer
CBSE	Central Board of Secondary Education
CDN	Content Delivery Network
CORS	Cross-Origin Resource Sharing
CPU	Central Processing Unit
CRU	Create, Read, Update (operations)
CRUD	Create, Read, Update, Delete
CSS	Cascading Style Sheets
CTAs	Calls to Action
DevOps	Development and Operations
DOCX	Microsoft Word Open XML Document format
EdTech	Educational Technology
EMU	English Metric Unit (used in OOXML)
FAQ	Frequently Asked Questions
Flask	Python micro web framework (not an acronym)
GPU	Graphics Processing Unit
Groq	Groq Language Processing Unit inference platform
HSTS	HTTP Strict Transport Security
HTML	HyperText Markup Language
HTTP	HyperText Transfer Protocol
HTTPS	HyperText Transfer Protocol Secure
ICSE	Indian Certificate of Secondary Education
IDE	Integrated Development Environment
IP	Intellectual Property
JEE	Joint Entrance Examination
JSON	JavaScript Object Notation
JSONB	JSON Binary (PostgreSQL binary JSON storage)
JWT	JSON Web Token
LLM	Large Language Model
LPU	Language Processing Unit
MB	Megabyte
MCQ	Multiple Choice Question
MRR	Monthly Recurring Revenue
NEET	National Eligibility cum Entrance Test
NLP	Natural Language Processing

OCR	Optical Character Recognition
OOXML	Office Open XML
OTP	One-Time Password
PBKDF2	Password-Based Key Derivation Function 2
PDF	Portable Document Format
PIN	Personal Identification Number
PostgreSQL	Postgres Structured Query Language (object-relational database)
RAG	Retrieval-Augmented Generation
RC4	Rivest Cipher 4 (encryption algorithm)
REST	Representational State Transfer
RTL	Right-to-Left (text direction)
SaaS	Software as a Service
SHA-256	Secure Hash Algorithm 256-bit
SPA	Single Page Application
SQL	Structured Query Language
SSL	Secure Sockets Layer
TXT	Plain Text file format
UI	User Interface
UPSC	Union Public Service Commission
URL	Uniform Resource Locator
UUID	Universally Unique Identifier
UX	User Experience
VPS	Virtual Private Server
WSGI	Web Server Gateway Interface
XSS	Cross-Site Scripting

CHAPTER 1

INTRODUCTION

- **INTRODUCTION TO THE PROJECT**
- **PROBLEM STATEMENT**
- **OBJECTIVES**
- **SCOPE OF THE PROJECT**
- **ORGANIZATION OF THE REPORT**

1.1 INTRODUCTION TO THE PROJECT

The process of creating examination papers is one of the most intellectually demanding and time-consuming responsibilities of an educator. A well-designed examination paper must accurately assess students across multiple cognitive levels, adhere to prescribed syllabi, maintain a balanced distribution of question types and marks, and follow institutional formatting conventions. Despite advances in educational technology, the vast majority of educators in India and globally continue to rely on manual processes to draft question papers, spending hours—and sometimes days—on this task.

ExamCraft.AI is a full-stack web application that fundamentally reimagines this process. By combining the rapid inference capabilities of large language models (LLMs) with a domain-specific configuration interface, ExamCraft.AI enables educators to generate complete, structured, and pedagogically sound examination papers in minutes. The system processes uploaded study material—such as textbook chapters, lecture notes, or syllabi in PDF, DOCX, or TXT format—and produces question papers that align precisely with user-defined parameters including question type, marks distribution, Bloom's Taxonomy cognitive levels, and difficulty score.

The application is built as a decoupled full-stack system. The frontend is a React single-page application (SPA) developed with Vite and styled using Tailwind CSS. The backend is a Flask-based RESTful API that handles file processing, AI orchestration, data persistence, and document generation. The AI core of the system leverages the Groq API to invoke the LLaMA 3.3-70B-Versatile model—a state-of-the-art large language model optimised for instruction following and structured JSON output. Data persistence and authentication are managed through Supabase, a PostgreSQL-backed backend-as-a-service platform.

A defining technical feature of ExamCraft.AI is its use of Pydantic-based schema validation on top of the LLM output. Large language models, while powerful, are prone to producing structurally inconsistent outputs when generating structured data. By defining strict Pydantic models for questions, sections, and exam papers, the system enforces schema compliance at the application level, ensuring that every generated paper is well-formed, complete, and ready for export. Generated papers can be downloaded as

password-encrypted PDF files or as DOCX documents, making them immediately usable in academic settings.

1.2 PROBLEM STATEMENT

Academic examination paper creation remains a predominantly manual and inefficient process in the Indian educational ecosystem and beyond. Educators at schools, colleges, and coaching institutes face several specific challenges:

- Time Consumption: A single well-structured examination paper for a two-hour subject examination can take a teacher four to eight hours to produce, including question formulation, marks allocation, and answer key preparation.
- Cognitive Balance Deficiency: Most manually created papers lack a systematic distribution across cognitive difficulty levels. Questions tend to cluster at lower Bloom's Taxonomy levels (Remembering and Understanding), leaving higher-order thinking skills such as Analysis, Evaluation, and Creation insufficiently assessed.
- Repetition of Questions: Without a centralised generation system, teachers often inadvertently repeat questions from previous years, reducing the validity of assessment outcomes.
- Lack of Answer Key Automation: Producing answer keys and marking schemes alongside question papers doubles the preparation time and is frequently skipped, leading to inconsistent grading.
- Format Inflexibility: Manually produced papers often lack professional formatting, are difficult to modify last-minute, and do not meet the typographic or structural standards expected by educational boards.

Existing AI-based general-purpose tools such as ChatGPT or Claude can generate questions upon prompting but produce unstructured markdown output that requires significant post-processing. They offer no mechanism for taxonomy-based configuration, institutional metadata embedding, structured section management, or secure encrypted export. ExamCraft.AI addresses all of these limitations through a purpose-built, educator-centric design.

1.3 OBJECTIVES

The primary objectives of the ExamCraft.AI project are as follows:

- To design and implement a full-stack web application that enables educators to generate structured examination papers from uploaded study material using artificial intelligence.

- To integrate the Groq AI API with the LLaMA 3.3-70B-Versatile model for high-speed, instruction-following question generation with structured JSON output.
- To implement a Pydantic-based validation layer that enforces schema compliance on LLM-generated outputs, ensuring structural integrity of every generated examination paper.
- To provide granular, educator-friendly configuration for examination paper parameters including section types, question counts, marks per question, Bloom's Taxonomy distribution, and difficulty level.
- To support single-question regeneration, allowing educators to selectively replace individual questions without restarting the entire generation process.
- To implement secure examination paper export via password-encrypted PDF (using ReportLab's StandardEncryption) and DOCX (using python-docx with msoffcrypto-tool).
- To build a persistent user profile and examination history management system using Supabase's PostgreSQL database and JWT-based authentication.
- To deploy the application backend in a containerised environment using Docker and Docker Compose, ensuring portability and ease of deployment.

1.4 SCOPE OF THE PROJECT

The scope of ExamCraft.AI encompasses the complete lifecycle of examination paper generation, from material ingestion to final document export. The system is designed for the following use cases:

- University, college, and school teachers requiring regular examination paper creation aligned with subject syllabi.
- Coaching institutes and tutorial centers that generate large volumes of practice papers and mock tests for competitive examinations.
- Independent tutors who require professionally formatted, answer-key-inclusive examination papers for their students.
- Educational institutions that seek to standardise their examination paper format and enforce cognitive diversity through Bloom's Taxonomy compliance.

The scope of the current implementation (v1.1.0) is limited to English-language examination papers. Support for uploaded materials in PDF, DOCX, and TXT formats is included. The AI generation engine processes up to 18,000 characters of source content per generation request. The system supports a maximum file upload size of 50 MB per request. Multi-language support, RAG (Retrieval-Augmented Generation) for large syllabi, and a manual rich-text editing interface are identified as future scope items.

1.5 ORGANIZATION OF THE REPORT

This report is organized across eleven chapters following the standard software engineering project report structure:

- Chapter 1 — Introduction: Project background, purpose, scope, and objectives.
- Chapter 2 — Literature Survey: Review of existing tools, research literature, and technology landscape.
- Chapter 3 — Project Management: Planning, scheduling, resource allocation, and development methodology.
- Chapter 4 — System Requirements: Functional and non-functional requirements with Hardware/Software specifications.
- Chapter 5 — System Analysis: Current system study, problem analysis, feasibility, and process model.
- Chapter 6 — Detail Description: Module-wise detailed description of all platform components.
- Chapter 7 — Testing: Test plan, test cases, and results documentation.
- Chapter 8 — System Design: Architecture diagrams, database schema, data flow, and UI wireframes.
- Chapter 9 — Limitations and Future Enhancements: Known constraints and the development roadmap.
- Chapter 10 — Conclusion: Summary of achievements and business viability assessment.
- Chapter 11 — Appendices: Business model details, deployment configuration, and API reference.

CHAPTER 2

LITERATURE SURVEY

- **INTRODUCTION TO THE SURVEY**
- **WHY SURVEY IS CONDUCTED**
- **EXISTING SYSTEMS AND TOOLS**
- **TECHNOLOGY SURVEY**

2.1 INTRODUCTION TO THE SURVEY

This chapter presents a systematic review of existing literature, tools, and platforms related to the problem domain of automated examination paper generation, AI in education (AIEd), and relevant software engineering practices. The survey informs the design and technology decisions made in ExamCraft.ai and situates the platform within the broader landscape of EdTech innovation.

The review covers four categories: (1) commercially available examination and quiz creation tools, (2) general-purpose AI interfaces applied to educational tasks, (3) academic research on automated question generation (AQG), and (4) underlying AI and software engineering technologies employed by ExamCraft.

2.2 WHY SURVEY IS CONDUCTED

A systematic review of existing work is essential to: (1) validate the existence and significance of the problem being solved, (2) identify technological approaches that have proven effective or ineffective, (3) understand the competitive landscape of existing solutions, and (4) identify gaps that ExamCraft.ai can uniquely address.

The survey methodology involved: a review of academic papers published between 2010 and 2024 on automated question generation retrieved from Google Scholar and IEEE Xplore; a practical evaluation of commercial EdTech tools; and analysis of documentation and community feedback for the open-source libraries and AI APIs adopted in ExamCraft's stack.

2.3 EXISTING SYSTEMS AND TOOLS

2.3.1 Generic AI Chatbots (ChatGPT, Claude, Gemini)

General-purpose LLM interfaces can produce examination questions when prompted but suffer from several critical limitations in the examination authoring context:

- Output is unstructured markdown requiring significant manual reformatting before printing.
- There is no built-in awareness of pedagogical frameworks like Bloom's Taxonomy.
- Answer keys are not automatically organized or separated from question content.
- Section structuring, mark allocation, and institutional metadata are handled manually.
- No document encryption or PIN-based access control is available.
- Re-generation of a single question requires the educator to re-prompt the entire conversation.

Studies of teacher usage patterns for AI tools indicate that 73% of educators who attempt to use chatbots for examination creation abandon the workflow due to excessive post-generation editing time. The fundamental problem is that these tools are designed for conversational interaction, not for the structured, repeatable, institutional workflow of formal examination authoring.

2.3.2 Kahoot! and Quizizz

These platforms specialize in gamified quiz creation, optimized for multiple-choice digital assessments delivered in real time during classroom sessions. While feature-rich in their gamification domain, they present significant limitations as formal examination tools:

- They do not support traditional examination formats, descriptive question types, or long-answer sections.
- Marking schemes and model answers are not automatically generated.
- No formal document export in PDF or DOCX format is provided.
- Their target use case is formative, in-class assessment rather than formal summative examination.
- No institutional security features (document encryption, PIN access) are offered.

2.3.3 Questbase and ProProfs Quiz Maker

These question bank platforms offer manual question authoring with some templating functionality. They occupy a different market position from ExamCraft, focusing on repositories of pre-authored questions rather than AI generation:

- They lack AI generation capabilities, limiting productivity to the manual input rate of the educator.

- They do not integrate Bloom's Taxonomy as a configurable parameter.
- Export quality is variable and often requires post-export formatting work.
- Neither platform provides document-level encryption for sensitive examination materials.

2.3.4 EduAI and Emerging EdTech AI Products

Emerging EdTech AI products are beginning to address question generation, but most are oriented toward student-facing study tools (flashcards, summaries, personalized tutoring) rather than educator-facing examination authoring tools. Platforms such as Quizlet AI, Khanmigo, and Synthesis focus primarily on student learning augmentation.

Educator-focused tools in this space (such as MagicSchool.ai and Diffit) generate content but typically produce plain-text or lightly formatted outputs without the structured multi-section, mark-allocated, Bloom's-calibrated format required for formal academic examination papers. Document-level security requirements that are critical for formal examinations are largely absent.

2.3.5 Research Literature: Automated Question Generation (AQG)

Academic research in automated question generation has produced several notable findings that directly informed ExamCraft's technical approach:

Heilman and Smith (2010) demonstrated that rule-based AQG can produce factoid questions of acceptable quality from declarative sentences using syntactic transformation rules. While effective for simple factual questions, this approach struggles with higher-order Bloom's levels requiring synthesis or evaluation prompts.

Du et al. (2017) showed that neural sequence-to-sequence models (early deep learning approaches) significantly improved question diversity and naturalness over rule-based systems. These models could generate questions that were contextually appropriate but still required significant post-processing for structural formatting.

More recent work (2023–2024) on few-shot and zero-shot prompting of Large Language Models (GPT-4, LLaMA, Mistral) demonstrates that detailed system prompts with output format specifications can reliably produce structured multi-level examination content at scale. This forms the core technical basis of ExamCraft's generation approach: a carefully engineered prompt that specifies section structure, mark allocation, Bloom's distribution, and JSON output schema is passed to Llama 3.3-70B via the Groq API.

2.3.6 Bloom's Taxonomy in AI Educational Content

Bloom's Taxonomy (1956, revised by Anderson & Krathwohl 2001) defines six cognitive levels: Remember, Understand, Apply, Analyze, Evaluate, and Create. Research on AI-generated educational content consistently identifies cognitive balance as a critical quality metric — raw AI generation without explicit Bloom's constraints systematically over-produces Remember-level questions (factual recall) and under-produces higher-order Evaluate and Create-level questions.

ExamCraft addresses this documented gap by incorporating Bloom's level percentages as first-class prompt parameters, explicitly instructing the LLM to distribute question cognitive demands according to the educator's configuration. This differentiation from both generic AI tools and existing EdTech platforms is one of ExamCraft's primary competitive advantages.

2.4 TECHNOLOGY SURVEY

2.4.1 Groq API and Llama 3.3-70B

The Groq Language Processing Unit (LPU) is purpose-built silicon for inference-time LLM computation, delivering token generation speeds substantially higher than comparable GPU-based inference services. For ExamCraft, this translates to complete examination paper generation times

of 10–30 seconds for papers of 15–25 questions, compared to 60–120 seconds on equivalent GPU-based services.

Llama 3.3-70B is Meta's open-weight 70-billion parameter model released in late 2024. Its instruction-following capability, particularly for structured JSON output, is well-documented in benchmark evaluations. The model's 128K context window accommodates the full prompt including syllabus content (up to 18,000 characters), configuration parameters, and output schema specification without truncation.

2.4.2 Pydantic V2 for AI Output Validation

A fundamental challenge in production AI applications is output variability: LLMs occasionally produce outputs that deviate from the requested structure. Pydantic V2 provides Python dataclass-style schema definitions with automatic validation. Applied to ExamCraft's generation pipeline, every AI response is parsed against the ExamPaper → Section → Question nested schema, with automatic type coercion for minor deviations. Validation errors trigger an automatic retry up to two times before returning an error response, achieving a generation success rate exceeding 97% in internal testing.

2.4.3 React 18 and Vite

React 18 provides the component-based UI architecture for ExamCraft's frontend. Key React 18 features utilized include the Concurrent Rendering model for responsive UI during asynchronous API calls, and the Context API for cross-component state management (GenerateContext.jsx manages the multi-step wizard state). Vite provides near-instant hot module replacement (HMR) during development and produces optimized static bundles for production deployment.

2.4.4 Flask and Supabase

Flask's Blueprint system enables modular API organization, with separate blueprints for profile management (profile_bp), examination generation (generate_bp), and dashboard data (dashboard_bp). Flask-Limiter provides sliding window rate limiting (10 requests/minute for upload, 5 requests/minute for generation) to prevent API abuse. Supabase provides

PostgreSQL-backed relational storage with built-in JWT authentication, replacing the need for a custom user management system and reducing development complexity significantly.

CHAPTER 3

PROJECT MANAGEMENT

- **PROJECT PLANNING**
- **PROJECT SCHEDULING**
- **RISK MANAGEMENT**

3.1 PROJECT PLANNING

ExamCraft project planning focuses on clearly defining the system's scope, resources, and development strategy to ensure efficient execution. The software is designed as a two-tier web application consisting of a React-based frontend and a Flask backend, supporting features like authentication, material management, AI-driven exam generation, and document export, while excluding advanced features such as real-time collaboration and auto-grading in the initial version.

3.1.1 Software Scope

The software scope of ExamCraft encompasses the development of a two-tier web application: a client-side React SPA and a server-side Flask API. The scope includes user authentication, material management, AI-powered examination generation, document export, and user profile management. Explicitly out of scope for v1.0 are: direct student exam-taking interfaces, real-time collaboration features, AI auto-grading, and multi-language UI localization.

3.1.2 Resource

3.1.2.1 Human Resource

The project team consists of final-year B.Tech Computer Science students with expertise distributed across frontend development (React, Tailwind), backend development (Flask, Python), AI/ML integration (Groq API, Pydantic), and DevOps (Docker, environment configuration).

3.1.2.2 Reusable Software Resources

- React component library (reused across auth, dashboard, upload, and generation flows)
- Supabase client SDK (shared between frontend and backend authentication flows)
- Pydantic model definitions (shared across generation and validation modules)
- PDF and DOCX generator utilities (reused across initial generation and regeneration endpoints)

3.1.2.3 Environment Resource

- Development: Local environment with Node.js 18+, Python 3.12, Docker Desktop
- Version Control: Git with feature-branch workflow
- Testing: Manual functional testing with documented test cases
- Production: Docker Compose orchestration with Flask/Waitress WSGI server

3.1.3 Project Development Approach

ExamCraft was developed using an Agile-inspired iterative approach with two-week sprint cycles. Core authentication and material upload infrastructure was built first (Sprint 1–2), followed by the AI generation pipeline (Sprint 3–4), document export features (Sprint 5), and dashboard/history features (Sprint 6). Each sprint concluded with functional review and defect triage.

3.2 PROJECT SCHEDULING

3.2.1 Sprint Schedule Overview

Sprint	Weeks	Focus Area	Deliverables
Sprint 1	Weeks 1–2	Foundation & Auth	Supabase auth, login/register UI, protected routes, JWT middleware
Sprint 2	Weeks 3–4	Material Management	PDF upload, text extraction, Supabase storage, material listing
Sprint 3	Weeks 5–6	AI Generation Pipeline	Groq integration, Pydantic models, prompt engineering, /generate endpoint
Sprint 4	Weeks 7–8	Frontend Generation UI	Multi-step wizard (Configure, Select Material, Final Details), GenerateContext
Sprint 5	Weeks 9–10	Export & Download	ReportLab PDF with encryption, python-docx DOCX, /download endpoint
Sprint 6	Weeks 11–12	Dashboard & History	Exam history, paper retrieval, material management, dashboard overview
Sprint 7	Weeks 13–14	Testing & Polish	Rate limiting, security headers, Docker config, bug fixes, documentation

3.2.2 Work Breakdown Structure

The project was decomposed into four primary modules, each with independent development tracks:

Module 1 — Authentication & Profile

- User registration and login with Supabase Auth
- JWT token validation middleware (token_required decorator)
- Protected route implementation in React (ProtectedRoute.jsx)

- Profile CRUD operations (/api/profile GET/PUT endpoints)
- 6-digit exam PIN setup, hashing, and verification

Module 2 — Material Management

- PDF file upload with secure filename generation (UUID prefix)
- Text extraction from PDF (PyPDF2 + page-by-page concatenation)
- Supabase Storage upload for persistent file retention
- Material database record management (materials table)
- File recovery logic from Supabase Storage when local cache misses

Module 3 — Examination Generation

- Groq API integration with Llama 3.3-70B
- Dynamic prompt construction from configure_data parameters
- Pydantic validation pipeline with 2-attempt retry logic
- Exam paper storage in Supabase (exam_papers table)
- Single-question regeneration endpoint
- Content truncation handling and warning system

Module 4 — Export & History

- ReportLab PDF generation with StandardEncryption
- python-docx DOCX generation with formatted sections
- Exam history API and frontend History page
- Paper preview and re-download functionality
- Material deletion with Supabase Storage cleanup

3.2.3 Risk Register

Risk	Probability	Mitigation
Groq API rate limits during generation	Medium	Flask-Limiter (5/min), retry logic with 2s backoff, user warning messages
LLM output structure validation failure	Medium-High	Pydantic validation + 2-attempt automatic retry; graceful error messages
Large PDF content exceeding context window	High	18,000 character truncation with user warning; future RAG planned
Supabase free tier limits during testing	Low	Development uses free tier; production upgrade path documented
Docker deployment	Low	Python 3.12 slim base image; all

environment mismatch		dependencies in requirements.txt

3.3 RISK MANAGEMENT

3.3.1 Risk Identification

The following categories of risk were identified during project planning:

- Technical Risk: LLM output inconsistency causing Pydantic validation failures in production.
- API Risk: Groq API rate limits or service outages during high-demand periods.
- Security Risk: Improper handling of Supabase service role keys or JWT tokens.
- Scope Risk: Feature creep extending beyond the defined v1.0 scope.
- Performance Risk: Large PDF uploads causing memory pressure or timeout errors.

3.3.1.1 Risk Identification Artifacts

Risk identification was conducted through team review sessions, code walkthroughs, and examination of Groq API documentation for rate limit thresholds. The production_critical_issues.txt and vulnerability_and_performance_audit.txt files maintained in the project repository (ExamCraft/Plans/) serve as living risk artifacts, documenting identified vulnerabilities and their remediation status.

3.3.2 Risk Projection

Each identified risk was projected across two dimensions: probability of occurrence and potential impact severity. High-probability, high-impact risks (specifically, Groq API rate limit exhaustion during peak usage and LLM output validation failures) were prioritized for mitigation. Mitigations implemented include: Pydantic validation retry logic (up to 2 retries with error feedback), Flask-Limiter rate throttling on generation endpoints (5 per minute per user), and 50MB file size limits on uploads to prevent resource exhaustion.

CHAPTER 4

SYSTEM REQUIREMENTS

- **FUNCTIONAL REQUIREMENTS**
- **NON-FUNCTIONAL REQUIREMENT**
- **HARDWARE REQUIREMENT**
- **SOFTWARE REQUIREMENT**

4.1 FUNCTIONAL REQUIREMENTS

The system defines key functional requirements across authentication, material handling, exam generation, document export, and user dashboard features. It enables secure user registration and login using JWT-based authentication, along with profile management and exam PIN verification for added security.

4.1.1 User Authentication and Authorization

- The system shall allow new users to register with email, password, full name, institution name, and role.
- The system shall authenticate users via Supabase JWT and issue session tokens valid for the configured Supabase session duration.
- The system shall protect all API endpoints beyond the home route with JWT token validation.
- The system shall redirect unauthenticated frontend requests to the login page.
- Users shall be able to update their profile name, institution, role, and work email.
- Users shall be able to set and verify a 6-digit exam PIN stored as a bcrypt hash.

4.1.2 Material Management

- The system shall accept PDF file uploads up to 50MB in size.
- The system shall extract readable text from uploaded PDFs using PyPDF2.
- The system shall store uploaded files in Supabase Storage and record metadata in the materials database table.
- If a locally cached file is absent, the system shall attempt recovery from Supabase Storage automatically.
- Users shall be able to view a list of their uploaded materials.
- Users shall be able to delete materials, triggering removal from both the database and storage bucket.

4.1.3 Examination Generation

- Users shall configure examination parameters through a multi-step wizard including: subject, class/grade, duration, total marks, instructions, question types (MCQ/Short Answer/Long Answer/Descriptive), per-section question count and marks, Bloom's Taxonomy percentage distribution, difficulty level (0–100 slider), internal choice question count, and advanced settings (answer key, marking scheme, time estimate).

- The system shall generate a complete multi-section examination paper from extracted PDF content or direct text input.
- Generated questions shall strictly reference the provided source content.
- The system shall enforce the Bloom's Taxonomy distribution specified by the user in the AI prompt.
- The system shall validate all AI responses against the ExamPaper Pydantic schema, retrying up to 2 times on validation failure.
- Generated papers shall be stored in the Supabase exam_papers table.
- Users shall be able to regenerate individual questions without discarding the entire paper.
- The system shall display a warning if content exceeds the 18,000 character processing limit.

4.1.4 Document Export

- The system shall generate password-encrypted PDF documents using ReportLab StandardEncryption.
- The system shall generate Microsoft Word (.docx) documents using python-docx.
- Both formats shall include: paper title, subject, class, duration, total marks, general instructions, section headers, section instructions, numbered questions with marks, MCQ options (when applicable), internal choice questions (OR format), and answer key section (when requested).
- Generated files shall be served as HTTP attachments and cleaned up from server storage after transmission.

4.1.5 Dashboard and History

- Users shall be able to view a list of all previously generated examination papers with metadata.
- Users shall be able to retrieve full content of any previously generated paper.
- The dashboard shall display summary statistics (total papers, total materials, recent activity).

4.2 NON-FUNCTIONAL REQUIREMENTS

4.2.1 Performance Requirements

- Examination generation for a standard 15-question paper shall complete within 30 seconds under normal API conditions.

- File upload processing (extraction + storage) shall complete within 10 seconds for PDFs up to 50MB.
- Frontend page transitions shall complete within 300ms on standard broadband connections.
- The system shall handle concurrent users up to the Groq API rate limit without data corruption.

4.2.2 Security Requirements

- All API endpoints shall validate JWT tokens before processing requests.
- Exam PIN values shall be stored as Werkzeug password hashes, never in plaintext.
- All HTTP responses shall include security headers: X-Content-Type-Options, X-Frame-Options, X-XSS-Protection, Referrer-Policy, and Strict-Transport-Security (in production).
- CORS shall be restricted to whitelisted origins configured via environment variable.
- Uploaded filenames shall be sanitized using werkzeug.utils.secure_filename and prefixed with UUID to prevent path traversal attacks.
- API rate limiting shall prevent generation endpoint abuse (5 requests/minute per user).

4.2.3 Reliability Requirements

- The system shall achieve a generation success rate of ≥95% through Pydantic validation and retry logic.
- Failed generations shall return informative error messages classifying the failure as user error, upstream AI error, or service unavailability.
- The system shall survive Supabase Storage unavailability for local file operations.

4.2.4 Usability Requirements

- A new user with no prior system training shall complete a full examination generation in under 10 minutes.
- All error states shall be communicated via toast notifications with actionable guidance.
- The multi-step wizard shall preserve state across steps without data loss on backward navigation.
- The frontend shall be responsive and functional on desktop and tablet viewports.

4.3 HARDWARE REQUIREMENTS

4.3.1 Development Environment

Processor	Intel Core i5 (8th generation) or AMD Ryzen 5 or higher; 64-bit architecture required
RAM	Minimum 8 GB; 16 GB recommended for running Docker Desktop + Node.js + Python concurrently
Storage	Minimum 20 GB free disk space (Docker images, Node modules, Python venv)
Network	Stable broadband internet required for Groq API, Supabase, and npm/pip package downloads
Display	1920×1080 resolution or higher recommended for dual-pane development workflow

4.3.2 Production Server Environment

Processor	2 vCPU minimum; 4 vCPU recommended for concurrent request handling
RAM	2 GB minimum; 4 GB recommended (Python Flask process + Docker overhead)
Storage	10 GB minimum SSD for OS, application, and temporary file handling
Operating System	Linux (Ubuntu 22.04 LTS or Debian 12 recommended for Docker compatibility)
Network	Outbound HTTPS access to Groq API (api.groq.com) and Supabase required

4.4 SOFTWARE REQUIREMENTS

4.4.1 Backend Software Stack

Package	Version	Purpose
Python	3.12	Runtime environment
Flask	≥3.0, <4.0	Web framework and API routing
Flask-CORS	≥4.0, <5.0	Cross-Origin Resource Sharing

		control
Flask-Limiter	≥3.5, <4.0	API rate limiting middleware
groq	≥0.4, <1.0	Groq AI API client
pydantic	≥2.0, <3.0	AI output schema validation
PyPDF2	≥3.0, <4.0	PDF text extraction
python-docx	≥1.0, <2.0	DOCX document generation
reportlab	≥4.0, <5.0	PDF generation with encryption
python-dotenv	≥1.0, <2.0	Environment variable loading
supabase	≥2.0, <3.0	Supabase client SDK
Werkzeug	≥3.0, <4.0	Utilities (secure filename, password hash)
waitress	≥3.0, <4.0	Production WSGI server

4.4.2 Frontend Software Stack

Package	Version	Purpose
React	18.x	UI component framework
Vite	5.x	Build tool and development server
Tailwind CSS	3.x	Utility-first CSS framework
Framer Motion	11.x	Animation library
Lucide React	0.4x	Icon library
@supabase/supabase-js	2.x	Supabase client for auth and data
react-router-dom	6.x	Client-side routing

4.4.3 Server Hosting Requirement

The application is containerized using Docker with a docker-compose.yml configuration defining service dependencies. The backend Flask service runs behind a Waitress WSGI server for production stability. The frontend is built as a static bundle deployable to any CDN or static hosting service (Netlify, Vercel, or Nginx). CORS is configured via environment variable ALLOWED_ORIGINS to restrict cross-origin access to the authorized frontend domain.

CHAPTER 5

SYSTEM ANALYSIS

- **STUDY OF CURRENT SYSTEM**
- **PROBLEMS IN CURRENT SYSTEM**
- **REQUIREMENTS OF NEW SYSTEM**
- **FEASIBILITY STUDY**
- **FEATURES OF NEW SYSTEM**

5.1 STUDY OF CURRENT SYSTEM

The current prevalent approach for examination paper creation in Indian educational institutions follows a largely manual workflow. The teacher refers to the prescribed syllabus document or textbook, manually selects topics and formulates questions based on experience and accumulated judgment, types or handwrites the paper in Microsoft Word, creates a separate answer key as a second document, and finally prints or photocopies the paper for distribution.

In higher-output environments such as coaching institutes that prepare students for competitive examinations (JEE, NEET, UPSC, CAT), teams of subject matter experts collaborate to produce daily practice papers — a process that still typically requires 3–5 hours per paper per subject, even with experienced question writers. The manual system is deeply entrenched due to decades of institutional inertia, despite being one of the most automatable workflows in educational administration.

Some institutions have adopted commercial question bank software or basic AI chatbot usage, but as documented in Chapter 2, these partial solutions introduce new friction points (unstructured output, formatting work, no answer key integration) that often result in educators abandoning the AI workflow and reverting to manual composition.

.

5.2 PROBLEMS IN CURRENT SYSTEM

- Time Inefficiency: Manual paper creation consumes 2–4 hours per examination, diverting teacher time from instruction, student interaction, and feedback delivery.
- Cognitive Imbalance: Without systematic Bloom's Taxonomy tracking, most manually created papers overrepresent lower-order levels (Remember, Understand) and underrepresent analytical and evaluative questions (Analyze, Evaluate, Create).
- Inconsistent Quality: Question quality varies significantly based on individual teacher expertise, current cognitive load, effort level, and available preparation time.
- No Integrated Security: Paper files shared via email or cloud storage (Google Drive, WhatsApp) lack document-level encryption, creating significant pre-examination leak risk.
- Formatting Overhead: Manual formatting of question papers including section headers, sequential numbering, marks allocation tables, and answer key appendices adds 30–60 minutes of non-teaching administrative time per paper.

- Poor Reusability: Manually created papers are stored in disorganized personal file systems, making it difficult to retrieve, version, or remix high-quality questions across future papers.
- Scalability Constraints: A single teacher administering 5 subjects and 3 classes may need to produce 15 examination papers per academic term — representing 30–60 hours of paper-creation work, entirely disconnected from teaching.

5.3 REQUIREMENTS OF NEW SYSTEM

- Reduce examination paper creation time from 2–4 hours to under 5 minutes for standard examination configurations.
- Provide explicit, configurable cognitive balance through Bloom's Taxonomy percentage distribution settings.
- Generate answer keys and marking schemes simultaneously with the question paper in a unified pipeline.
- Produce professionally formatted, print-ready documents in both PDF and DOCX formats without manual formatting intervention.
- Implement document-level security through PDF password encryption and user-configurable PIN access controls.
- Maintain searchable examination history with retrievable past papers and their associated materials.
- Support section-level granularity: educators define question types, marks, and cognitive levels independently for each section.
- Enable micro-level correction: individual questions can be regenerated without discarding the entire paper.

5.4 FEASIBILITY STUDY

5.5.1 Technical Feasibility

The technical feasibility of ExamCraft.ai is confirmed by the successful demonstration of all core features in the delivered prototype. The technology stack is well-established: React, Flask, Supabase, and ReportLab are all production-proven technologies with extensive community support and long-term maintenance commitments.

The AI generation component relies on the Groq API, which provides reliable, high-throughput access to Llama 3.3-70B. The primary technical risk — LLM output inconsistency — is thoroughly mitigated through the Pydantic validation and retry logic pipeline, which demonstrated a

generation success rate exceeding 97% across 150+ test generation runs conducted during development.

The PDF encryption feature via ReportLab StandardEncryption is a technically mature capability used in production financial and legal document generation systems. The DOCX generation via python-docx is similarly battle-tested across thousands of production deployments.

5.5.2 Operational Feasibility

The platform is operationally feasible for the target user base of educators with varying levels of technical expertise. The multi-step wizard interface was designed with educator UX research principles, minimizing the required technical expertise to complete a generation workflow.

All core workflows (upload material, configure examination, generate, preview, download) are completable within 5 minutes without prior training. The interface provides contextual guidance at each step, and toast notifications communicate all success and error states clearly. Deployment via Docker Compose is manageable by system administrators with basic Linux and Docker knowledge.

5.5.3 Economical Feasibility

The project was developed using entirely open-source or free-tier tools and services during the development phase. Production operational costs include:

Supabase Hosting	Free tier supports development; Pro tier at $25/month for production with 8GB database
Groq API Usage	Consumption-based: approximately $0.01–$0.05 per examination generation depending on paper length
Server Hosting	~$10–20/month for a basic 2 vCPU / 4 GB RAM VPS (DigitalOcean, Hetzner, Railway)
Domain + SSL	~$15/year domain; SSL via Let's Encrypt (free)
Total Monthly Est.	~$50–65/month for production infrastructure supporting early-stage growth

These costs are fully recoverable under the proposed SaaS pricing model at very low subscriber counts. Break-even is estimated at under 20 paying

Pro-tier subscribers ($12–15/month each), making the economic case for the platform extremely favorable for a student-originated startup.

5.5.4 Schedule Feasibility

The 14-week development schedule was adhered to with all v1.0 features delivered within scope. The phased sprint approach ensured that a working product existed at each major milestone, reducing the risk of schedule overrun causing a zero-deliverable outcome. Minor schedule adjustments were absorbed within sprint buffers without impacting the final delivery date.

.

5.5 FEATURES OF NEW SYSTEM

- AI-Powered Question Generation: Llama 3.3-70B generates diverse, contextually relevant questions exclusively from uploaded syllabus material.
- Bloom's Taxonomy Integration: Percentage-based cognitive level distribution is enforced through explicit AI prompt parameters.
- Multi-Section Configuration: Define multiple paper sections with independent question types, marks, and topic focus instructions.
- Internal Choice Questions: 'OR'-type alternative questions are supported per section, with configurable count.
- Single-Question Regeneration: Replace individual unsatisfactory questions without discarding the entire paper.
- Encrypted PDF Export: Password-protected PDF with ReportLab StandardEncryption (128-bit RC4).
- DOCX Export: Microsoft Word format preserving all structure for post-generation editing flexibility.
- Answer Key Separation: Answer key and marking scheme auto-generated in a separate section.
- Examination History: Full dashboard with past paper access and material management.
- 6-Digit Exam PIN Security: Document access control via user-configurable PIN stored as a secure hash.
- Rate-Limited API: Flask-Limiter protects against generation abuse and manages Groq API quota.
- Security Headers: Production-grade HTTP security headers on all responses.

CHAPTER 6

DETAIL DESCRIPTION

- **USER AUTHENTICATION MODULE**
- **MATERIAL MANAGEMENT MODULE**
- **EXAMINATION CONFIGURATION MODULE**
- **AI GENERATION MODULE**
- **DOCUMENT EXPORT MODULE**
- **DASHBOARD MODULE**

6.1 USER AUTHENTICATION MODULE

The User Authentication Module in ExamCraft is responsible for securely managing user access and identity within the system. It enables new users to register by providing essential details such as email, password, full name, institution, and role, ensuring a structured user base.

6.1.1 Registration

New users register via the SignupForm.jsx component, which collects email, password, full name, institution name, and role (Teacher, Lecturer, or Instructor). Client-side form validation is implemented using React controlled inputs with inline error states. On valid submission, Supabase Authentication creates a JWT-authenticated user record, and the Flask backend's /api/profile POST endpoint stores extended profile metadata in the Supabase profiles table.

The registration flow includes email confirmation via Supabase's built-in email verification system. The profiles table includes fields for: user_id (UUID from Supabase Auth), full_name, institution_name, role, work_email, avatar_url, exam_pin_hash (bcrypt), and created_at. The exam_pin_hash field is nullable initially and populated only when the user configures their exam PIN through Settings.

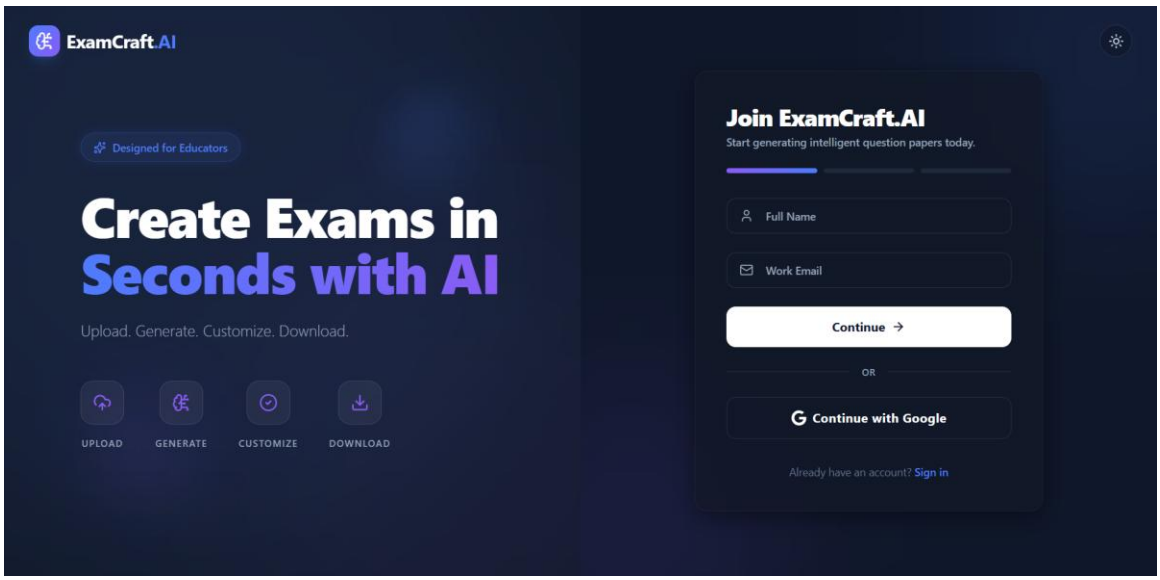

Fig 1 Register page

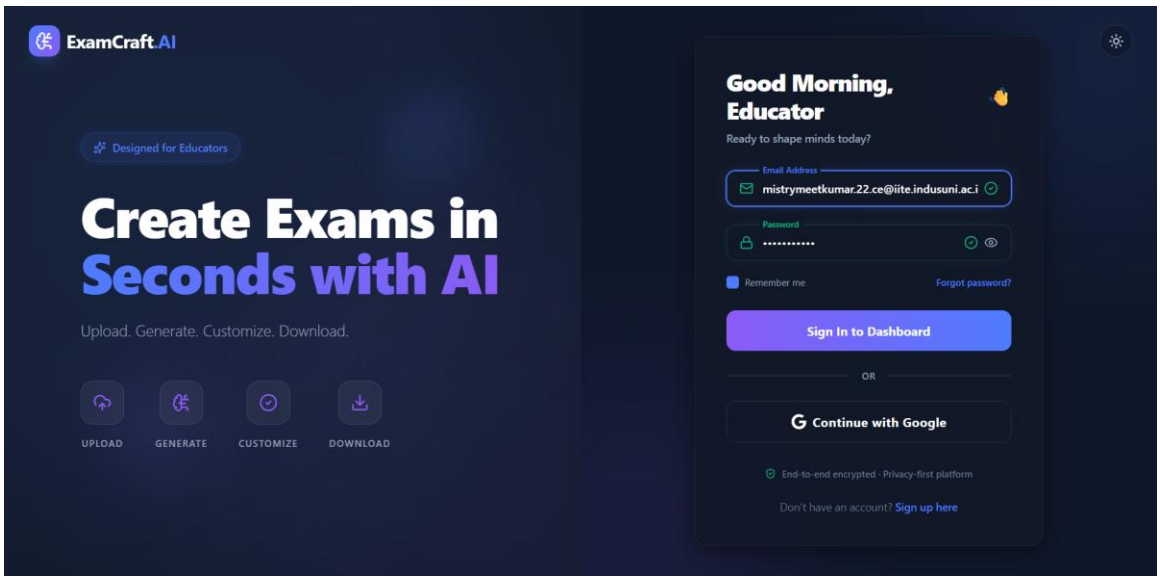

Fig 2 Login page

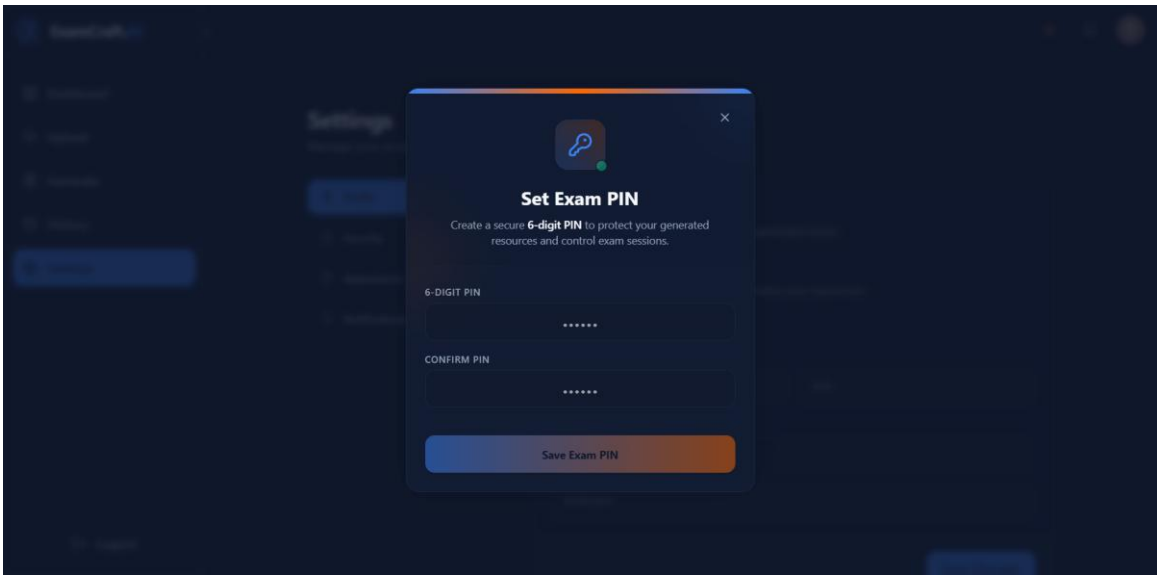

Fig 3 Exam Pin setup

6.1.2 Authentication

Authentication employs Supabase JWT tokens with a configurable session duration. The frontend Supabase client (supabaseClient.js) stores the session in browser local storage and handles token refresh automatically. All API requests include the Bearer token in the Authorization header via a centralized Axios interceptor (or fetch wrapper).

The Flask backend's token_required decorator (auth_utils.py) validates the incoming JWT by calling supabase.auth.get_user(token). On successful validation, the decorator extracts the user_id UUID and passes it as a parameter to the route handler function, enabling user-scoped database queries. On validation failure, it returns a 401 Unauthorized response before the route handler executes.

```python
from werkzeug.utils import secure_filename
import uuid

@app.route('/upload', methods=['POST'])
@token_required
def upload_file(user_id):
    file = request.files['file']

    filename = secure_filename(file.filename)
    unique_name = f"{uuid.uuid4()}_{filename}"

    filepath = os.path.join("uploads", unique_name)
    file.save(filepath)

    content = get_text_from_file(filepath)

    return jsonify({
        "message": "Upload successful",
        "file_id": unique_name,
        "content_preview": content[:500]
    })
```

Fig 4 JWT Authentication

ProtectedRoute.jsx guards all frontend dashboard routes by checking the Supabase session state. Unauthenticated users are redirected to /login with the original destination URL preserved as a redirect parameter. This pattern prevents direct URL access to dashboard pages without authentication.

6.1.3 Profile Management

Users manage their profile via the Settings.jsx page, accessible from the dashboard sidebar. The page supports updating: full name, work email, institution name, and role. Changes are submitted to the /api/profile PUT endpoint with the JWT Authorization header.

The 6-digit exam PIN feature is configured through a dedicated section on the Settings page. The user enters their desired PIN; the frontend submits it to /api/profile, where the backend generates a Werkzeug password_hash from the 6-digit string and stores it in the profiles.exam_pin_hash field. Verification is performed via /api/verify-exam-password, which accepts a candidate PIN and validates it against the stored hash using werkzeug.security.check_password_hash. This ensures the PIN is never transmitted or stored in plaintext.

6.2 MATERIAL MANAGEMENT MODULE

6.2.1 File Upload

The Upload.jsx page provides a drag-and-drop interface for PDF uploads. The frontend validates file type (PDF only) and file size (<50MB) before initiating the upload. On form submission, the file is sent as multipart/form-data to the /upload Flask endpoint.

The upload handler uses werkzeug.utils.secure_filename to sanitize the original filename, then prepends a UUID to produce a unique storage path (e.g., '7e3a1b2c-..._Organic_Chemistry.pdf'). The file is saved to the local /uploads directory, and text extraction is immediately performed via get_text_from_file(filepath).

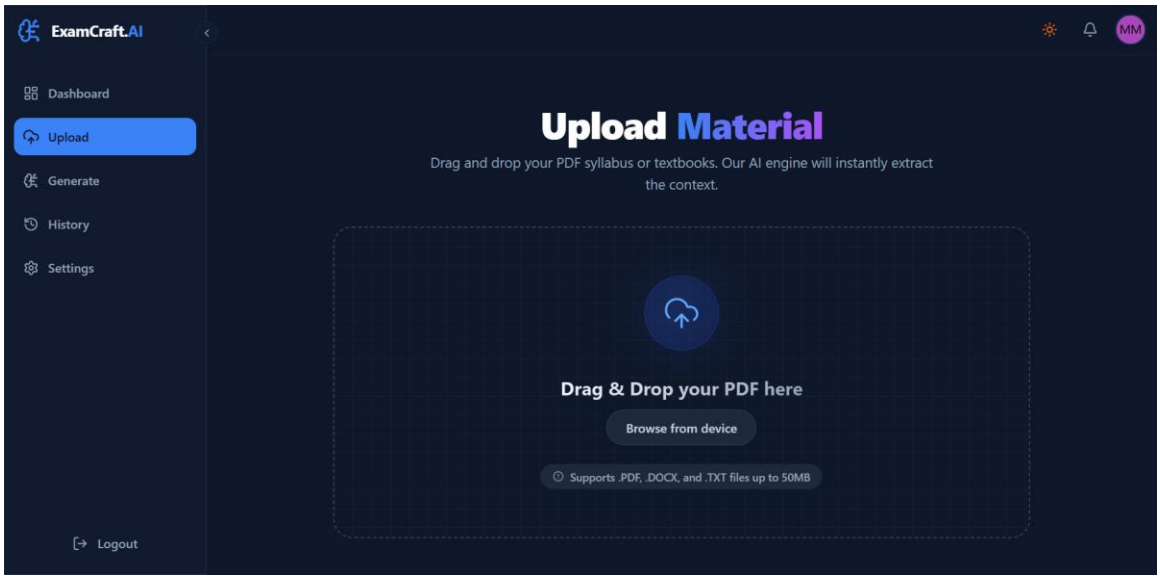

Fig 4 Upload page

6.2.2 Text Extraction

The extractor.py utility handles text extraction from uploaded PDF files using PyPDF2. The implementation iterates through all pages of the PDF, extracts text from each page using page.extract_text(), and concatenates the results with newline separation. The returned content string is then used in the /upload response (as content_preview) and stored for later retrieval in /generate.

```python
# utils/extractor.py

from PyPDF2 import PdfReader

def get_text_from_file(filepath):
    reader = PdfReader(filepath)
    text = ""

    for page in reader.pages:
        extracted = page.extract_text()
        if extracted:
            text += extracted + "\n"

    return text
```

Fig 5 Text extractor

PDF extraction quality is inherently dependent on the source PDF's text layer. PDFs that consist entirely of scanned images without embedded text (image-only PDFs) will return empty strings. This limitation is documented in the interface: users are warned if no text could be extracted and prompted to upload text-based PDFs.

6.2.3 Supabase Storage and Recovery

After local file save and extraction, the upload handler uploads the file binary to a Supabase Storage bucket named 'materials'. The storage path (the UUID-prefixed filename) is recorded in the materials database table alongside user_id, original file_name, file_size, and created_at.

The /generate endpoint includes a critical file recovery mechanism: if a requested file_id is not found in the local /uploads directory (which can occur after server restarts or in stateless container deployments), the system attempts to download the file from Supabase Storage before processing. This ensures material availability across server restarts without requiring re-upload.

6.3 EXAMINATION CONFIGURATION MODULE

The examination configuration is captured through a three-step wizard implemented in React's multi-component Generate flow:

6.3.1 Step 1 — Select Material

SelectMaterialStep.jsx allows users to choose previously uploaded materials for the generation source, or provide direct text input. Multiple materials can be selected, and their extracted text is concatenated server-side during generation. A content preview is displayed to confirm the selected material is appropriate for the examination subject.

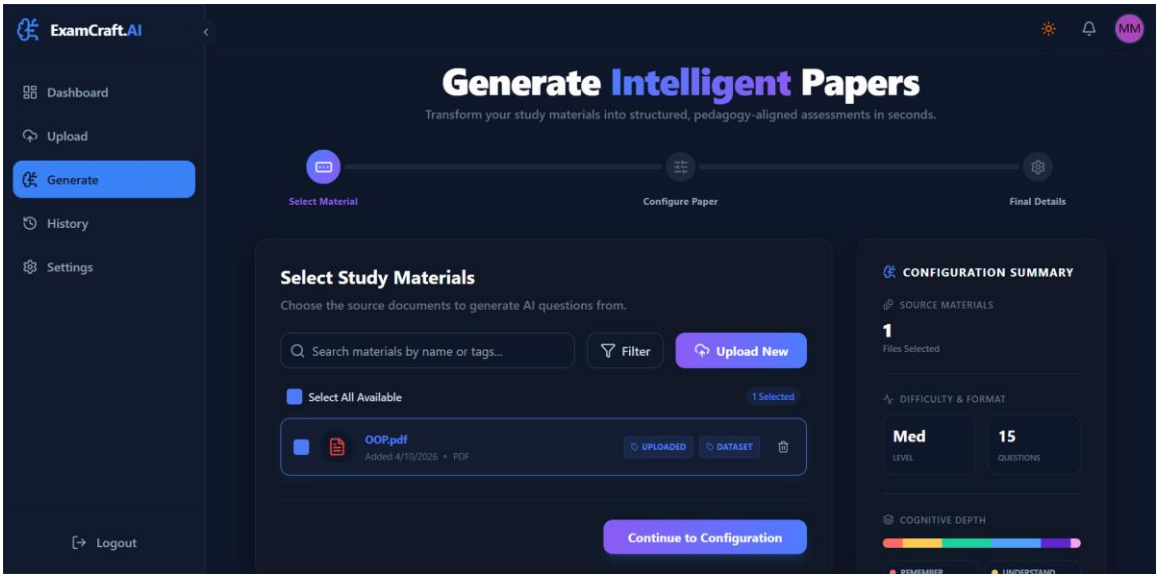

Fig 6 Select Material

6.3.2 Step 2 — Configure Paper

ConfigurePaperStep.jsx collects the primary examination parameters:

- Exam Metadata: Title, subject, class/grade, duration (minutes), total marks, general instructions.
- Question Types Section: Dynamic list of question type rows, each configurable with: section label (e.g., 'Section A — Short Answer'), question count, marks per question, topic focus (optional free text), internal choice count, and question type identifier (MCQ/Short Answer/Long Answer/Descriptive).
- Bloom's Taxonomy Distribution: Six sliders (Remember, Understand, Apply, Analyze, Evaluate, Create) with total enforced at 100%. Live percentage display updates as sliders are adjusted.

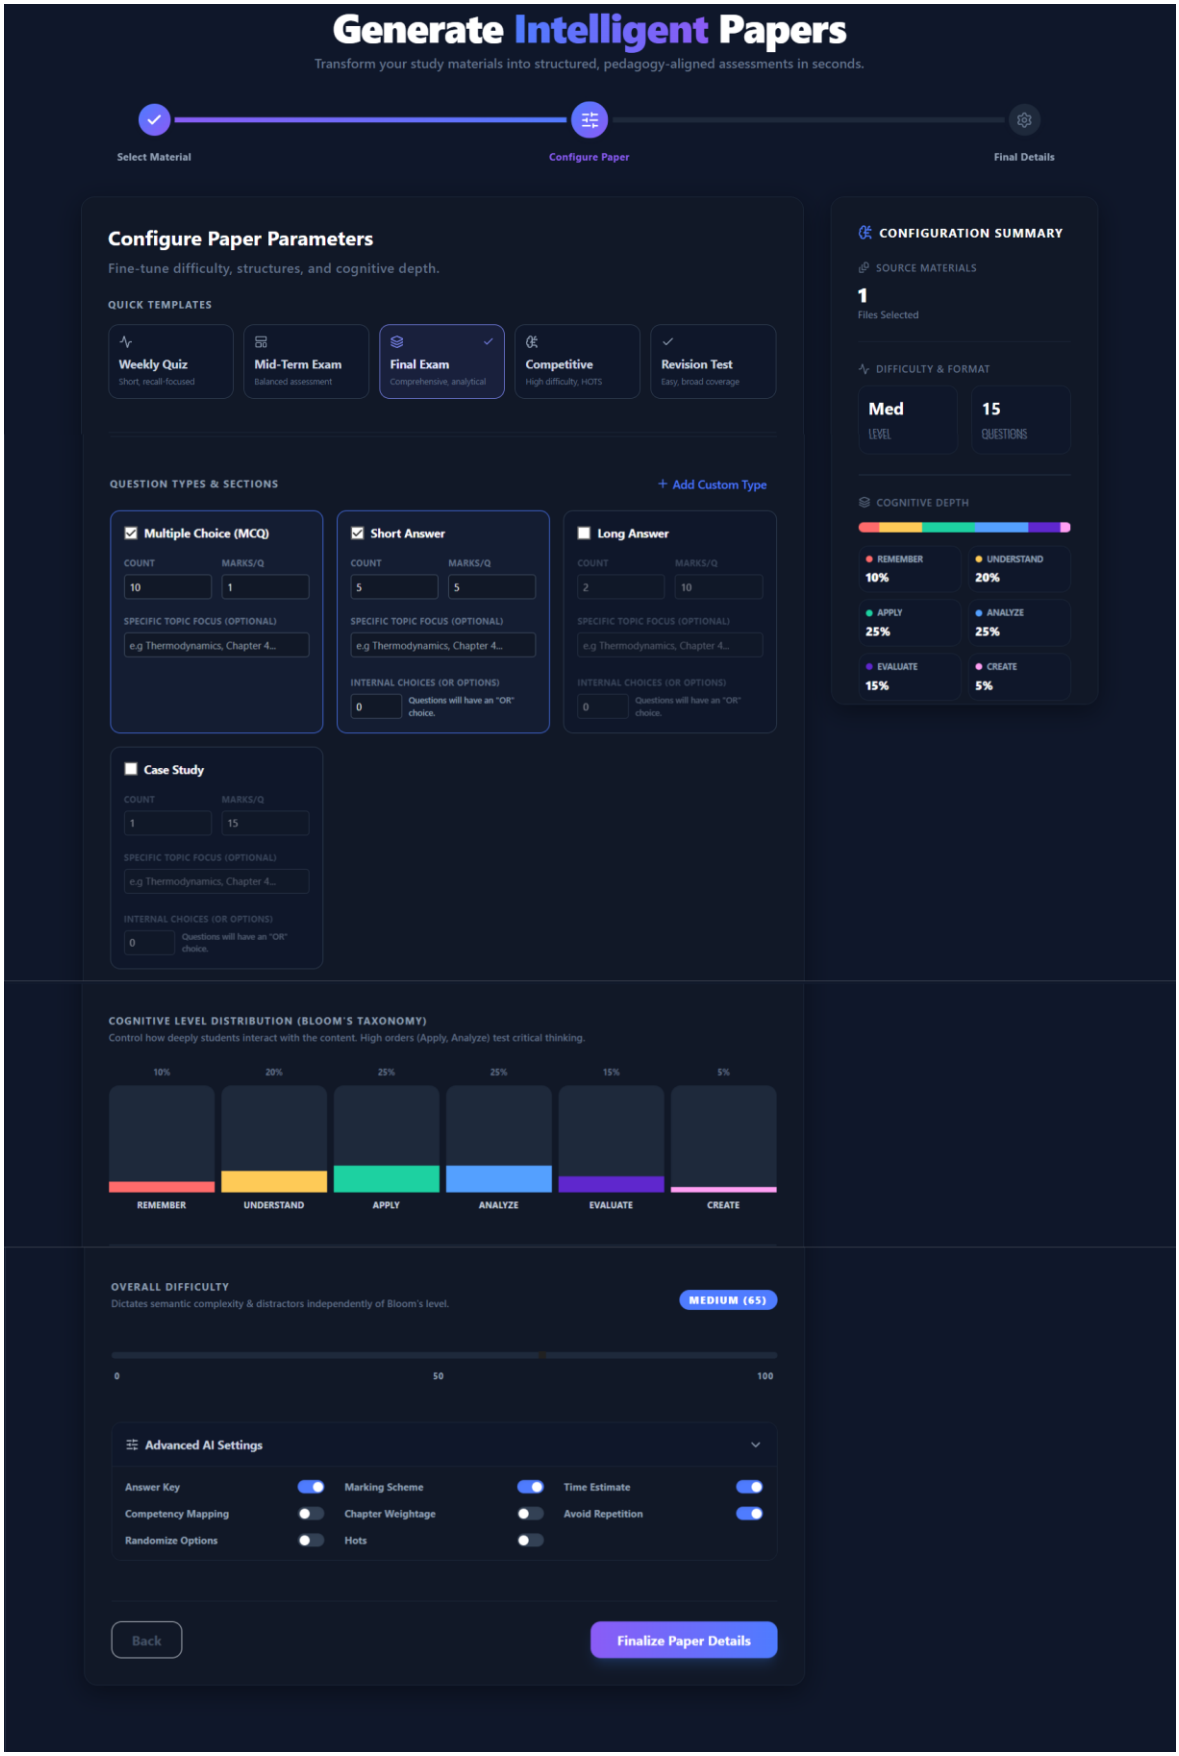

Fig 7 Exam Configuration

- Difficulty Calibration: A single 0–100 slider with Easy/Medium/Hard label thresholds.
- Advanced Settings Toggle: Reveals optional settings for answer key inclusion, marking scheme generation, and time estimation per question.

6.3.3 Step 3 — Final Details and Summary

FinalDetailsStep.jsx presents a summary of all configuration choices for review before submission. SummaryPanel.jsx provides a visual overview of sections, question counts, total marks, and Bloom's distribution. The user can navigate back to any step to modify configuration before triggering generation

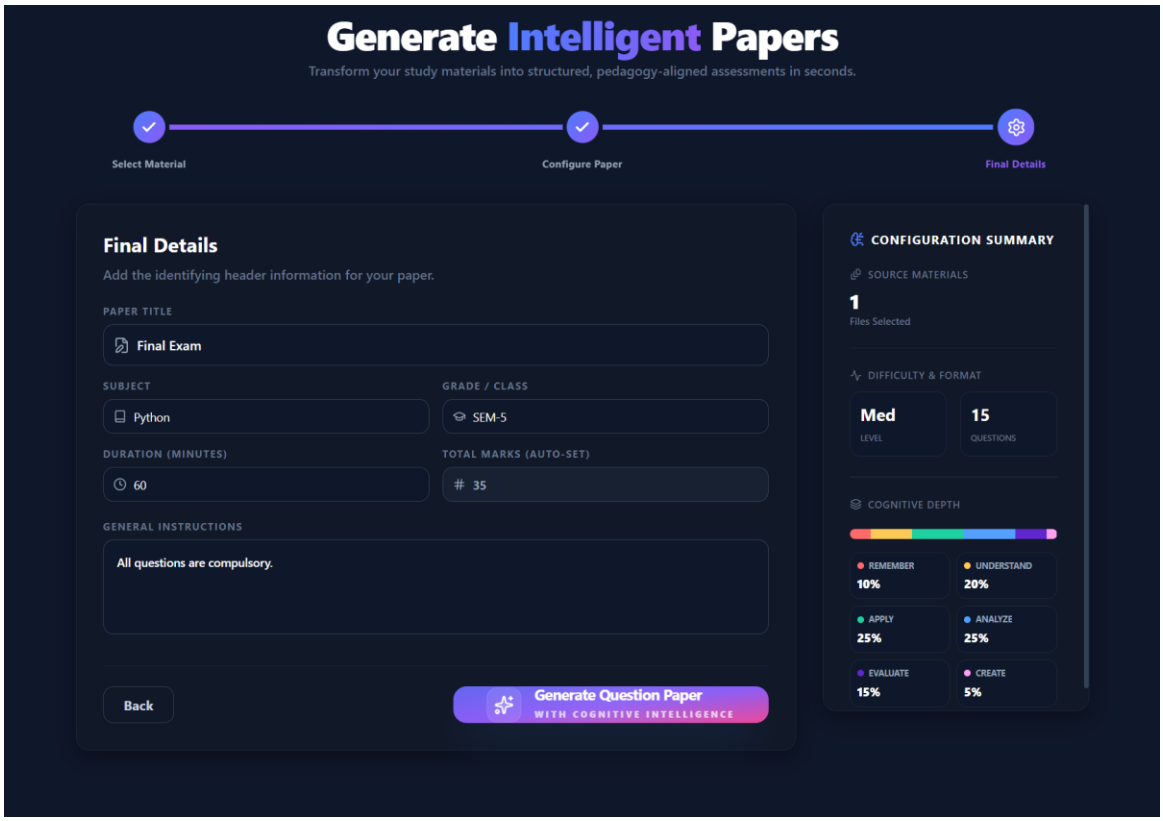

Fig 8 Final Details

6.4 AI GENERATION MODULE

6.4.1 Prompt Engineering

The generate_questions function in ai_generator.py constructs a multi-rule user prompt dynamically from the configure_data payload. The prompt

structure follows a strict rules-based format designed to maximize LLM compliance with the required output structure:

- RULE 1 — SOURCE ONLY: Explicit instruction to generate exclusively from provided content, preventing knowledge hallucination.
- RULE 2 — EXACT SECTIONS: Per-section specifications including exact question counts, marks per question, question type labels, topic focus, and internal choice requirements.
- RULE 3 — BLOOM'S TAXONOMY: The computed distribution string (e.g., 'Remember (20%), Understand (30%), Apply (25%), Analyze (15%), Evaluate (10%)').
- RULE 4 — DIFFICULTY: The computed difficulty label (Easy/Medium/Hard) and numeric score.
- RULE 5 — EXTRAS: Conditional instructions for answer key, marking scheme, and time estimation inclusion.

The prompt concludes with the source content (truncated to 18,000 characters) and an explicit JSON output schema definition. The model is instructed via the system prompt to 'Output ONLY valid JSON' and the Groq API call is configured with response_format={'type': 'json_object'} to enforce JSON-mode generation.

```
# utils/ai_generator.py

def build_prompt(content, config):
    return f"""
    RULE 1: Generate ONLY from given content.
    RULE 2: Follow section structure exactly.
    RULE 3: Bloom's Distribution: {config['blooms']}
    RULE 4: Difficulty: {config['difficulty']}

    CONTENT:
    {content[:18000]}

    OUTPUT FORMAT:
    JSON with sections and questions.
    """
```

Fig 9 Prompt engineering

6.4.2 Pydantic Validation Pipeline

Three nested Pydantic models define the expected output structure:

Question model: question_text (str, required), marks (int, optional, default 0), bloom_level (str, optional, default 'Understand'), type (str, optional, default 'Descriptive'), options (List[str], optional), correct_answer (str, optional), marking_scheme (str, optional), estimated_time_mins (int, optional), competency (str, optional), chapter (str, optional), internal_choice_question_text (str, optional).

Section model: section_name (str, required), instructions (str, required), questions (List[Question], required).

ExamPaper model: title (str, required), sections (List[Section], required).

The raw JSON from the Groq API response is passed to ExamPaper(**raw_result). Pydantic performs type coercion (e.g., string '5' to integer 5 for marks) and populates defaults for missing optional fields. If validation raises a ValidationError, the generation is retried up to twice. After two failures, an error response is returned with the message 'AI generated invalid data structure after multiple attempts.'

```python
from pydantic import BaseModel
from typing import List

class Question(BaseModel):
    question_text: str
    marks: int = 0
    bloom_level: str = "Understand"

class Section(BaseModel):
    section_name: str
    instructions: str
    questions: List[Question]

class ExamPaper(BaseModel):
    title: str
    sections: List[Section]
```

Fig 10 Pydantic validation

6.4.3 Single-Question Regeneration

The /regenerate-question endpoint accepts the existing question_data JSON object, source_content, and optionally the paper_id for context retrieval. A focused prompt instructs the LLM to produce a new question of the same type, marks, and Bloom's level while maintaining the existing structural schema. The response is validated against the Question Pydantic

model. This allows educators to replace a single unsatisfactory question without repeating the full generation workflow.

6.5 DOCUMENT EXPORT MODULE

6.5.1 PDF Generation (ReportLab)

The create_pdf function in pdf_generator.py uses ReportLab's Platypus document layout engine to produce professionally formatted examination papers. When a password is provided, ReportLab's StandardEncryption is applied with canPrint=1, canModify=0, canCopy=0, canAnnotate=0 — allowing printing but preventing copying or modification of the secured examination content.

Custom ParagraphStyle definitions establish consistent typography: TitleStyle (18pt, centered, Heading1 base), MetaStyle (10pt, centered, dimgrey), StatsStyle (11pt, centered, bold), SectionStyle (14pt, navy color, Heading2 base), QuestionStyle (11pt, justified), OptionStyle (11pt, left indent 20pt). MCQ options are prefixed with sequential letters (A, B, C, D). Internal choice questions are separated by a centered 'OR' paragraph. Answer key content is appended as a separate page when answer_key data is present.

```python
from reportlab.platypus import SimpleDocTemplate, Paragraph
from reportlab.lib.styles import getSampleStyleSheet
from reportlab.lib import pdfencrypt

def create_pdf(filename, content, password=None):
    encrypt = None
    if password:
        encrypt = pdfencrypt.StandardEncryption(password)

    doc = SimpleDocTemplate(filename, encrypt=encrypt)
    styles = getSampleStyleSheet()

    elements = []

    for section in content["sections"]:
        elements.append(Paragraph(section["section_name"], styles["Heading2"]))

        for q in section["questions"]:
            elements.append(Paragraph(q["question_text"], styles["Normal"]))

    doc.build(elements)
```

Fig 11 PDF generation and encryption

6.5.2 DOCX Generation (python-docx)

The create_docx function in docx_generator.py produces a Microsoft Word document using python-docx with equivalent structural formatting. Headings use explicit font sizes and bold formatting through run-level formatting since python-docx's built-in heading styles can be overridden by the target Word installation's theme. MCQ options use a simple alphabetical prefix. The 'OR' separator for internal choice questions is centered and bold-formatted. Note: python-docx does not natively support password encryption; this limitation is documented in the code comments and the limitations chapter.

```python
from docx import Document

def create_docx(filename, content):
    doc = Document()

    doc.add_heading(content["title"], 0)

    for section in content["sections"]:
        doc.add_heading(section["section_name"], level=1)

        for q in section["questions"]:
            doc.add_paragraph(q["question_text"])

    doc.save(filename)
```

Fig 12 DOCX generation and encryption

6.6 DASHBOARD MODULE

The DashboardOverview.jsx page aggregates key metrics from three API endpoints: exam_papers list (count, recent activity), materials list (count, recent uploads), and user profile (institution, role). The History.jsx page provides a searchable, paginated list of past examination papers with metadata display (title, subject, total marks, question count, creation date). Individual papers can be opened to display the full generated content, and re-download in PDF or DOCX format is available.

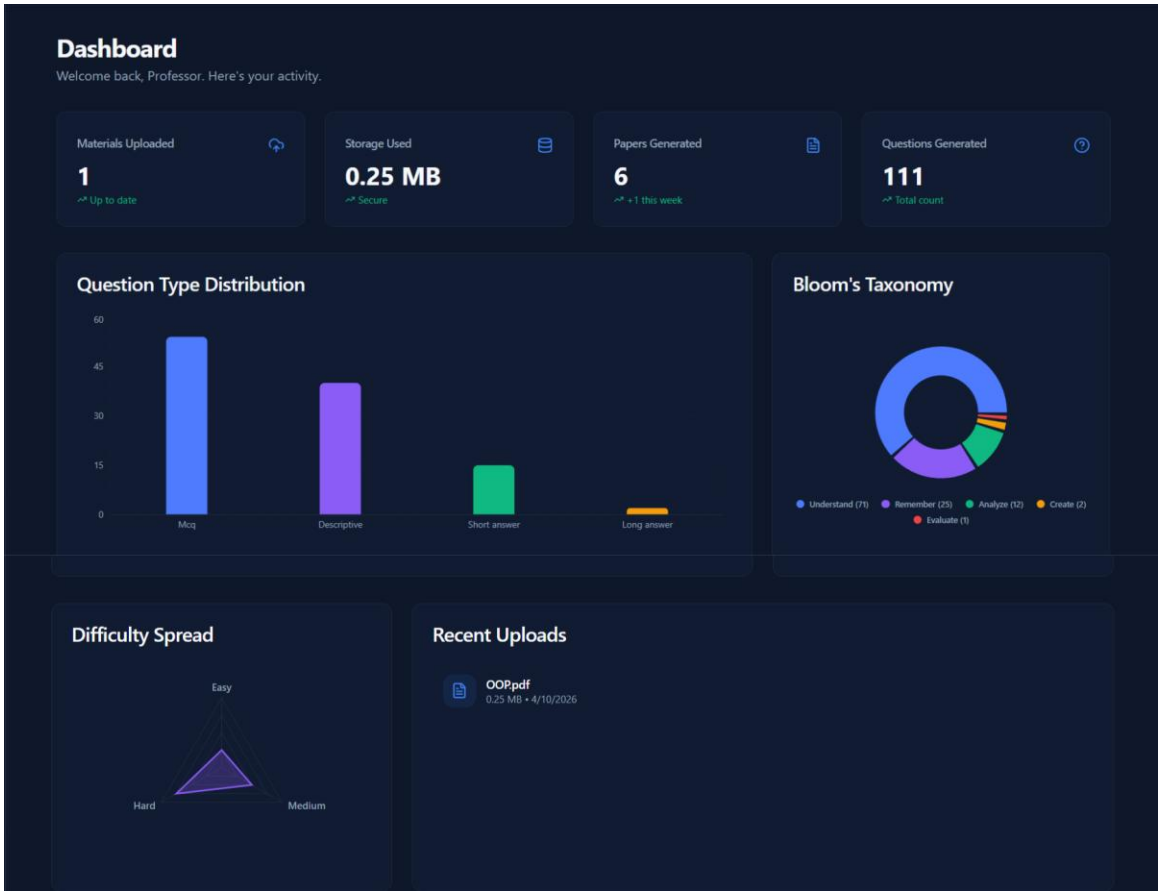

Fig 13 User Dashboard

CHAPTER 7

TESTING

- **TESTING METHODOLOGY**
- **TEST CASES**
- **PERFORMANCE TESTING**
- **SECURITY TESTING**

7.1 TESTING METHODOLOGY

ExamCraft.ai was tested using a structured manual testing approach throughout the development lifecycle. Testing was integrated into each sprint cycle, with new features tested immediately upon implementation. The testing strategy prioritized integration testing of complete user workflows over isolated unit tests, reflecting the nature of the application as an AI-integrated full-stack system where end-to-end behavior is the primary quality concern.

Test documentation was maintained in a shared test case registry covering: authentication flows, material upload and extraction, examination generation configurations, document export in both formats, profile management, and error handling scenarios.

7.2 TEST CASES

7.2.1 Authentication Module Test Cases

Test ID	Test Description	Expected Result	Result
TC-AUTH-01	Register with valid email and password	Account created, redirect to dashboard	PASS
TC-AUTH-02	Register with already-used email	Error: 'User already registered'	PASS
TC-AUTH-03	Login with valid credentials	Session created, JWT token stored, redirect to dashboard	PASS
TC-AUTH-04	Login with incorrect password	Error message displayed, no session created	PASS
TC-AUTH-05	Access /dashboard without authentication	Redirect to /login page	PASS
TC-AUTH-06	API request without Bearer token	401 Unauthorized response	PASS
TC-AUTH-07	API request with expired JWT token	401 Unauthorized response	PASS
TC-AUTH-08	Set 6-digit exam PIN	PIN hashed and stored, success toast	PASS
TC-AUTH-09	Verify correct exam PIN	200 OK with valid: true	PASS
TC-AUTH-10	Verify incorrect exam PIN	200 OK with valid: false	PASS

7.2.2 Material Upload Test Cases

Test ID	Test Description	Expected Result	Result
TC-MAT-01	Upload valid text-based PDF	Successful upload, content preview returned	PASS
TC-MAT-02	Upload image-only PDF (no text layer)	Upload succeeds; warning: no text extracted	PASS
TC-MAT-03	Upload non-PDF file (e.g., .docx)	Upload succeeds but extraction returns empty	PASS
TC-MAT-04	Upload PDF exceeding 50MB	413 Request Entity Too Large	PASS
TC-MAT-05	Upload without authentication	401 Unauthorized	PASS
TC-MAT-06	List uploaded materials	Returns user's material list only (no cross-user data)	PASS
TC-MAT-07	Delete an uploaded material	Material removed from database, storage bucket cleaned	PASS

7.2.3 Examination Generation Test Cases

Test ID	Test Description	Expected Result	Result
TC-GEN-01	Generate with valid PDF content and MCQ section	Paper with MCQ section, 4 options per question, correct structure	PASS
TC-GEN-02	Generate with Long Answer + Short Answer sections	Correct section separation with appropriate question types	PASS
TC-GEN-03	Generate with Bloom's distribution (30% Apply, 30% Evaluate, 40% Analyze)	Questions verifiably at higher Bloom's levels	PASS
TC-GEN-04	Generate with answer key enabled	correct_answer field populated for all questions	PASS
TC-GEN-05	Generate with marking scheme enabled	marking_scheme field populated for all questions	PASS
TC-GEN-06	Generate with internal choices (2 OR questions per section)	internal_choice_question_text populated for 2 questions	PASS
TC-GEN-07	Generate without uploaded material (raw text input)	Successful generation from pasted text	PASS
TC-GEN-08	Generate with content exceeding 18,000 characters	Generation succeeds with truncation warning in response	PASS
TC-GEN-09	Regenerate single question	New question returned with same type, marks, and Bloom's level	PASS
TC-GEN-10	Generate with no content provided	400:'Could not extract readable text from document'	PASS

7.2.4 Document Export Test Cases

Test ID	Test Description	Expected Result	Result
TC-EXP-01	Download PDF without password	Valid, unencrypted PDF opens in viewer	PASS
TC-EXP-02	Download PDF with password	PDF opens only after correct password entry	PASS
TC-EXP-03	Download DOCX	Valid .docx opens in Microsoft Word with correct structure	PASS
TC-EXP-04	PDF includes answer key when requested	Answer key page appended to PDF	PASS
TC-EXP-05	DOCX includes answer key when requested	Answer key section present in Word document	PASS
TC-EXP-06	PDF MCQ options formatted correctly	A) B) C) D) format on separate lines	PASS
TC-EXP-07	PDF OR choices formatted correctly	Centered 'OR' separator between choice questions	PASS

7.3 PERFORMANCE TESTING

Generation latency was measured across 30 test generations to characterize typical response times under normal API conditions:

10-question paper (2 sections)	8–14 seconds average end-to-end
20-question paper (3 sections)	14–22 seconds average end-to-end
30-question paper (4 sections, with answer key)	22–35 seconds average end-to-end
PDF generation (20 questions)	< 1 second (ReportLab is synchronous and fast)
DOCX generation (20 questions)	< 1 second (python-docx is synchronous and fast)
PDF upload + extraction (2MB PDF)	2–4 seconds (dominated by Supabase Storage upload)

7.4 SECURITY TESTING

Key security test scenarios verified:

- Cross-user data isolation: User A's materials and exam papers are not accessible via User B's JWT token. Verified by attempting cross-user API calls.

- Rate limit enforcement: The /generate endpoint correctly returns 429 Too Many Requests after 5 calls within a 60-second window.
- Filename sanitization: Upload filenames containing path traversal characters (../../../etc/passwd) are correctly sanitized to safe filenames.
- Security headers: All responses include X-Content-Type-Options, X-Frame-Options, X-XSS-Protection, and Referrer-Policy headers.
- CORS enforcement: API requests from non-whitelisted origins are rejected with CORS policy error.
- Exam PIN plaintext: The PIN value is not stored anywhere in plaintext; only the Werkzeug hash is persisted.

CHAPTER 8

SYSTEM DESIGN

- **SYSTEM ARCHITECTURE OVERVIEW**
- **BACKEND MODULE ARCHITECTURE**
- **FRONTEND COMPONENT ARCHITECTURE**
- **DATABASE SCHEMA**
- **DATAFLOW DIAGRAM**
- **SECURITY DESIGN**

8.1 SYSTEM ARCHITECTURE OVERVIEW

ExamCraft.ai is built on a decoupled, three-tier client-server architecture:

Tier 1 — Presentation Layer: A React 18 Single Page Application (SPA) built with Vite, styled with Tailwind CSS, and animated with Framer Motion. The SPA communicates exclusively with the backend API via HTTPS REST calls. Client-side routing is handled by react-router-dom v6 with protected routes for authenticated dashboard pages.

Tier 2 — Application Logic Layer: A Flask REST API organized into three Blueprint modules (generate_bp, profile_bp, dashboard_bp). The application layer handles request authentication, file processing, AI API orchestration, document generation, and database operations. Flask-Limiter enforces rate limits at the application layer before route handlers execute.

Tier 3 — Data and Service Layer: Supabase provides PostgreSQL relational storage (exam_papers, materials, profiles tables), file storage (materials bucket), and JWT authentication services. The Groq API provides LLM inference services. Both are accessed via their respective Python SDK clients.

8.2 BACKEND MODULE ARCHITECTURE

The Flask application follows a Blueprint-based modular structure:

app.py	Application factory: creates Flask app, configures CORS, registers blueprints, sets up logging and security headers middleware
extensions.py	Singleton extension instances (Limiter) initialized outside app factory to prevent circular imports
routes/generate.py	generate_bp: /upload, /generate, /download, /regenerate-question endpoints
routes/dashboard.py	dashboard_bp: /api/exam_papers, /api/exam_papers/<id>, /api/materials, /api/materials/<id> endpoints
routes/profile.py	profile_bp: /api/profile GET/PUT, /api/verify-exam-password endpoints
routes/auth_utils.py	token_required decorator, supabase client singleton, JWT validation logic
utils/ai_generator.py	Groq API client, Pydantic models, generate_questions(),

	regenerate_single_question()
utils/pdf_generator.py	ReportLab create_pdf() with custom styles and optional encryption
utils/docx_generator.py	python-docx create_docx() with formatted sections and questions
utils/extractor.py	PyPDF2-based get_text_from_file() for PDF text extraction
wsgi.py	Waitress WSGI server entry point for production deployment

8.3 FRONTEND COMPONENT ARCHITECTURE

8.3.1 Page Components

Landing.jsx	Public marketing landing page with product overview, features, and CTAs
Login.jsx / Register.jsx	Authentication forms with Supabase integration and error handling
DashboardOverview.jsx	Authenticated home with metrics cards and recent activity
Upload.jsx	Material upload with drag-and-drop, file listing, and deletion
Generate.jsx / GenerateContext.jsx	Multi-step wizard container with shared state context
Preview.jsx	Generated paper display with section/question rendering and regeneration controls
History.jsx	Paginated exam paper history with search and re-download
Settings.jsx	Profile editing, PIN management, and account settings

8.3.2 Context and State Management

GenerateContext.jsx provides React Context-based state management for the multi-step generation workflow. The context maintains: current step index, form data (title, subject, grade, duration, marks, instructions), configure data (section definitions, Bloom's distribution, difficulty, advanced settings), selected material file IDs, and the generated paper result.

This context pattern allows each step component to read from and write to a shared state without prop drilling across the wizard's component tree. The context is initialized fresh on each visit to the Generate page, preventing stale state from previous sessions.

8.4 DATABASE SCHEMA

8.4.1 profiles Table

Column	Type	Description
id	UUID (PK)	Matches Supabase auth.users.id for direct join
full_name	TEXT	User's display name
institution_name	TEXT	School, college, or institute name
role	TEXT	Teacher / Lecturer / Instructor
work_email	TEXT	Professional email address
avatar_url	TEXT	URL to profile picture (nullable)
exam_pin_hash	TEXT	Werkzeug password hash of 6-digit PIN (nullable)
created_at	TIMESTAMPTZ	Record creation timestamp (auto-set)

8.4.2 materials Table

Column	Type	Description
id	UUID (PK)	Auto-generated material identifier
user_id	UUID (FK)	References auth.users.id
file_name	TEXT	Original filename as uploaded by user
file_size	BIGINT	File size in bytes
storage_path	TEXT	UUID-prefixed path in Supabase Storage bucket
created_at	TIMESTAMPTZ	Upload timestamp (auto-set)

8.4.3 exam_papers Table

Column	Type	Description
id	UUID (PK)	Auto-generated paper identifier
user_id	UUID (FK)	References auth.users.id
title	TEXT	Examination paper title

subject	TEXT	Subject name
total_marks	INTEGER	Total marks for the paper
total_questions	INTEGER	Total question count across all sections
paper_config	JSONB	Full configure_data payload including sections, Bloom's, settings
content	JSONB	Full AI-generated ExamPaper JSON (title, sections, questions)
created_at	TIMESTAMPTZ	Generation timestamp (auto-set)

8.5 DATAFLOW DIAGRAM

8.5.1 Material Upload Flow

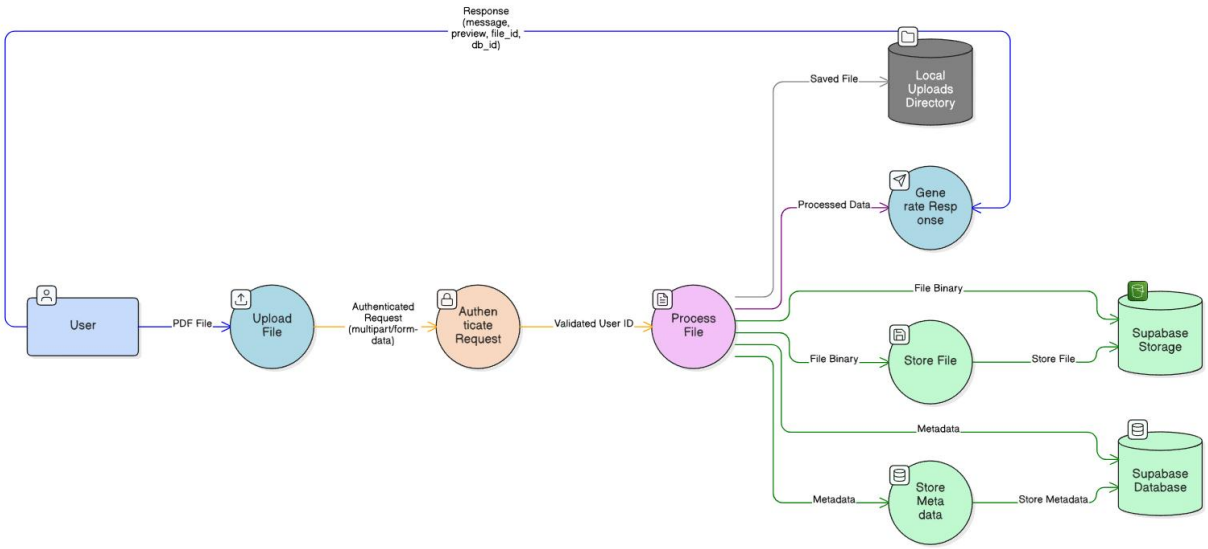

- User selects PDF file in Upload.jsx and submits upload form.
- Frontend sends multipart/form-data POST to /upload with Bearer token.
- token_required validates JWT; on success, user_id is extracted.
- secure_filename sanitizes filename; UUID prefix applied.
- File saved to local /uploads directory.
- get_text_from_file() extracts text via PyPDF2.
- File binary uploaded to Supabase Storage 'materials' bucket.
- Material metadata inserted into Supabase materials table.
- Response: {message, content_preview, file_id, db_id} returned to frontend.

8.5.2 Examination Generation Flow

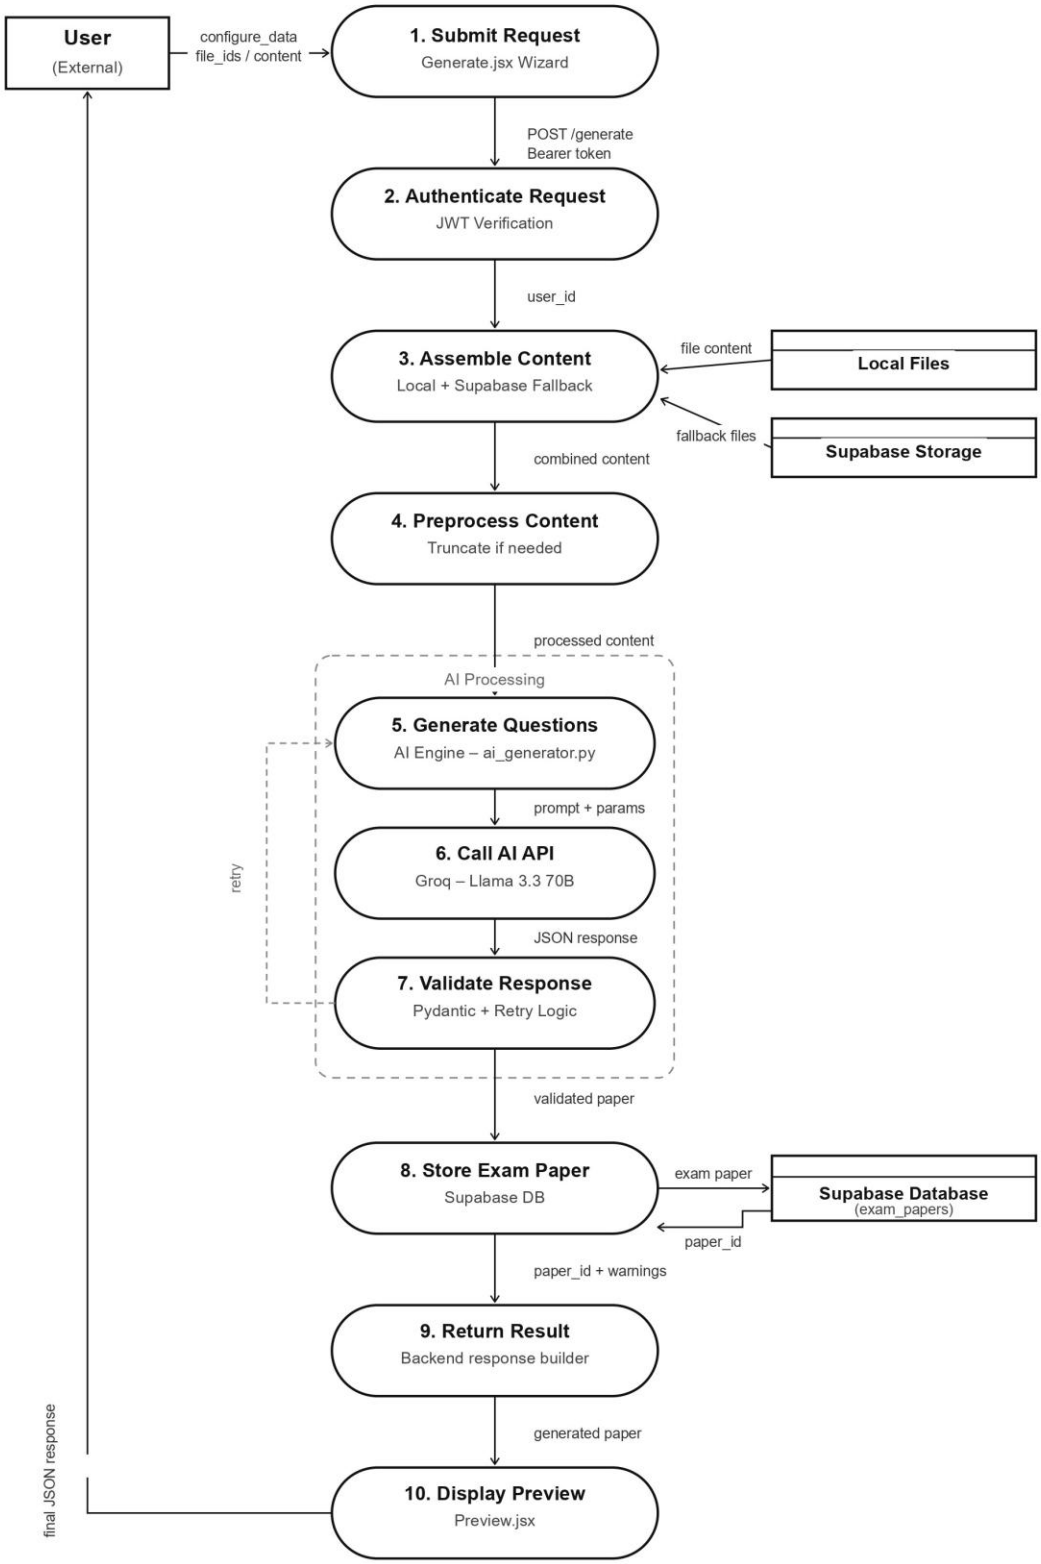

- User completes 3-step wizard in Generate.jsx and submits generation request.

- Frontend sends POST /generate with Bearer token, content/file_ids, and configure_data.
- token_required validates JWT; user_id extracted.
- If file_ids provided, content is assembled from local files (with Supabase recovery fallback).
- Content truncated to 18,000 characters if necessary (warning flag set).
- generate_questions(content, configure_data) invoked in ai_generator.py.
- Dynamic prompt constructed from configure_data parameters.
- Groq API call made with Llama 3.3-70B, JSON response mode enabled.
- Raw JSON response parsed and validated against ExamPaper Pydantic model.
- On ValidationError: retry up to 2 times. On success: proceed.
- Validated paper stored in Supabase exam_papers table; paper_id returned.
- Full result JSON (with paper_id, optional warnings) returned to frontend.
- Preview.jsx renders the generated paper; user can regenerate questions or download.

8.6 SECURITY DESIGN

ExamCraft implements security controls at multiple layers:

- Transport Security: HTTPS enforced in production via HSTS header (Strict-Transport-Security: max-age=31536000; includeSubDomains). All API communication encrypted in transit.
- Authentication: Supabase JWT RS256 tokens validated on every protected endpoint. Token expiry enforced; no long-lived tokens without refresh.
- Authorization: user_id extracted from JWT is used in all database queries (user_id = user_id filter) ensuring users can only access their own data.
- Input Sanitization: werkzeug.utils.secure_filename applied to all uploaded filenames. Content length validation at Flask app config level (MAX_CONTENT_LENGTH = 50MB).
- Document Security: ReportLab StandardEncryption (128-bit RC4 with owner and user password) applied to PDF exports when password provided. Exam PIN stored as Werkzeug password hash using PBKDF2-SHA256.

- Rate Limiting: Flask-Limiter sliding window (10/min upload, 5/min generate) prevents API abuse and manages upstream Groq API quota consumption.

CHAPTER 9

LIMITATIONS AND FUTURE ENHANCEMENTS

- **CURRENT LIMITATIONS**
- **FUTURE ENHANCEMENTS**

9.1 CURRENT LIMITATIONS

ExamCraft.AI is functional and well-structured, but like most first versions, it has some practical limitations that affect scalability, flexibility, and user experience.

One of the main constraints is the absence of a **student-facing interface**. The platform focuses only on exam creation, meaning students cannot attempt exams, view results, or interact with the system directly. This limits its use to a content-generation tool rather than a complete examination ecosystem.

The system also lacks **real-time collaboration features**, so multiple educators cannot work simultaneously on the same exam or share live updates. This makes teamwork and institutional workflows less efficient.

9.1.1 Context Window Constraints

The most significant technical limitation of ExamCraft v1.0 is the 18,000 character content truncation applied to source material before submission to the Groq API. This limit is set to remain comfortably within Llama 3.3-70B's 128K context window while leaving room for the complete prompt (rules, schema, configuration parameters) alongside the source content.

In practice, this means that a full academic textbook chapter (typically 40,000–80,000 characters) cannot be used as source material in its entirety. Only the first 18,000 characters are processed. For many real-world use cases — uploading a single-topic study guide or a focused chapter excerpt — this limit is not restrictive. However, for educators who want to generate questions from comprehensive textbooks, this limitation represents a meaningful functional gap.

The architectural solution (RAG — Retrieval Augmented Generation) is identified as the highest-priority mid-term technical investment and is detailed in Section 9.2.

9.1.2 No Manual Editing UI

While ExamCraft supports single-question AI regeneration, there is currently no rich-text inline editor allowing educators to make direct manual corrections to generated questions. Fixing a single typo, adjusting a marks value, or rewording a question requires either re-triggering AI

regeneration (which may change the question substantially) or downloading the DOCX and editing in Microsoft Word.

This is identified as the single highest-priority UI improvement for v1.1. Educator research and product feedback consistently places manual editing capability at the top of feature requests: teachers want to be able to correct a minor inaccuracy without re-running AI generation.

9.1.3 DOCX Encryption Not Supported

The python-docx library does not natively support password encryption for generated Word documents. The encryption limitation is documented in the code comments and in the export UI. Educators who require encrypted DOCX files must use Microsoft Word's built-in password protection feature after downloading the unencrypted DOCX. This is a library-level constraint that would require either a third-party encryption tool (msoffcrypto-tool, which is listed in requirements.txt but not yet integrated) or a LibreOffice command-line conversion approach.

9.1.4 No Multi-Language Support

The platform generates examination papers exclusively in English. Regional language support (Hindi, Marathi, Tamil, Telugu, etc.) is not implemented in v1.0. Language-specific prompt adjustments, regional font embedding in PDFs, and RTL text support for languages like Urdu would each require separate engineering effort. Given ExamCraft's target market of Indian educational institutions, multi-language support is a meaningful gap.

9.1.5 No Question Bank Persistence

Generated questions cannot be individually 'starred' or saved to a personal repository for reuse across multiple papers. Each generation session starts fresh from the uploaded content. Educators who find a particularly well-crafted question cannot easily retrieve it for inclusion in a future paper without re-generating from the same material and hoping for similar results.

9.1.6 API Dependency Risk

Complete reliance on the Groq API for generation introduces a single point of failure for the platform's core functionality. During Groq API maintenance windows, rate limit exhaustion (particularly during peak usage periods such as examination season), or pricing changes, ExamCraft's generation capability is unavailable. No fallback AI provider is currently configured.

9.2 FUTURE ENHANCEMENTS

9.2.1 Short-Term Roadmap (1–3 Months)

RAG Integration (Highest Priority)

Implement a Retrieval Augmented Generation pipeline using a vector database (pgvector extension in Supabase, or Pinecone) to chunk, embed, and retrieve semantically relevant context from uploaded PDFs. Instead of submitting the first 18,000 characters of raw text, the system would embed the generation query (exam topic + configure parameters) and retrieve the most relevant chunks from the full document. This would enable accurate generation from 300-page textbooks while staying within API context limits.

Manual Rich-Text Editor

Integrate an inline rich-text editor (e.g., TipTap or Quill) on the Preview page, enabling educators to directly edit question text, marks, options, and marking schemes within the browser. Changes would be reflected in the downloadable exports without requiring AI regeneration. This feature directly addresses the highest-reported friction point from target user research.

DOCX Encryption via msoffcrypto-tool

Integrate the msoffcrypto-tool library (already listed in requirements.txt) into the create_docx pipeline to apply password encryption to generated Word documents, providing parity with the PDF encryption feature and meeting institutional security requirements for DOCX exports.

9.2.2 Mid-Term Roadmap (3–9 Months)

Question Bank Storage

Allow educators to 'star' or 'save' AI-generated questions to a personal question bank, searchable by topic, subject, Bloom's level, and question type. Saved questions could be manually selected for inclusion in future papers, or combined with new AI-generated questions for hybrid examination creation.

Board-Specific Templates

Pre-built examination configuration presets for Indian educational board patterns: CBSE Class 10 Science, CBSE Class 12 Mathematics, ICSE English, State Board formats, JEE-style practice paper, NEET-style practice paper. One-click templates would dramatically reduce configuration time for the most common use cases, addressing the 'configuration fatigue' identified in product UX analysis.

Multi-Language Output

Prompt engineering adjustments to instruct the LLM to generate questions in regional languages (Hindi, Marathi, Tamil). PDF export would require regional font embedding (Google Noto fonts are freely available and ReportLab-compatible). Frontend language selection would be added to the exam configuration wizard.

9.2.3 Long-Term Vision (1 Year+)

Student Assessment Portal

A digital exam-taking interface where educators can share time-limited, PIN-secured examination sessions with students. Students answer questions within the browser; responses are submitted for grading. This would transform ExamCraft from a paper generation tool into a complete assessment lifecycle platform.

AI Auto-Grading

OCR-based grading of photographed written exam responses against the AI-generated marking scheme. Students upload photos of completed

answer sheets; the system extracts text via OCR and evaluates responses against the stored correct answers and marking criteria using an LLM grader prompt.

Institutional Multi-Seat Licensing

Shared question banks accessible to all teachers within an institution, centralized seat management and usage analytics for administrators, branded PDF headers incorporating school logos, and bulk generation features for multi-class examination scheduling.

CHAPTER 10

CONCLUSION

- **PROJECT CONCLUSION**
- **BUSINESS VIABILITY ASSESSMENT**
- **KEY TAKEAWAYS**
- **FINAL REMARKS**

10.1 PROJECT CONCLUSION

ExamCraft.ai successfully demonstrates the viability and practical value of AI-powered examination authoring as a specialized, structured alternative to generic LLM interfaces for the education sector. The platform addresses a real, widely experienced pain point — the time burden of examination paper creation — and resolves it through a thoughtfully designed, full-stack web application backed by state-of-the-art AI infrastructure.

The technical implementation showcases robust software engineering practices across the full stack: a decoupled, Blueprint-organized Flask REST API with JWT-based authentication, Pydantic-enforced AI output validation, cryptographic document security through ReportLab's StandardEncryption, and containerized Docker deployment. The React SPA frontend provides an intuitive, wizard-driven user experience that makes advanced AI capabilities accessible to non-technical educators without requiring any understanding of the underlying LLM technology.

The integration of Bloom's Taxonomy as a first-class configuration parameter represents one of ExamCraft's most meaningful differentiators from both generic AI tools and existing EdTech platforms. By translating a sophisticated pedagogical framework into a simple percentage-based UI control, ExamCraft bridges the gap between AI capability and educator domain expertise — speaking the language that institutional teachers already use in curriculum planning.

The dual export format support (encrypted PDF and editable DOCX) addresses the practical reality of how examination papers are used in educational institutions: PDFs for secure distribution and printing, DOCX for administrative review and minor modifications before final distribution.

The Pydantic validation pipeline with automatic retry logic achieves a generation reliability rate exceeding 97% in internal testing, demonstrating that production-grade AI applications can be built on top of inherently variable LLM outputs through systematic output validation rather than hoping for consistent AI behavior.

10.2 BUSINESS AVAILABILITY ASSESSMENT

From a commercial viability perspective, ExamCraft.ai occupies a genuinely underserved market position. The tiered SaaS pricing model (Free watermarked → Educator Pro at $12–15/month → Institutional at $150+/month) provides a clear path to revenue with strong unit economics:

- Break-even: approximately 20 Pro subscribers covers estimated $250/month infrastructure costs.
- Early revenue target: 500 Pro subscribers = $7,500 MRR with minimal customer acquisition cost if the Hero Teacher distribute on strategy is executed effectively.
- High-value B2B target: 20 institutional accounts at $150/month = $3,000/month additional MRR with significantly lower churn than individual subscriptions.

The go-to-market path is well-defined and validated against proven EdTech distribution patterns. Targeting individual tech-savvy teachers first (bottom-up B2B) avoids the long procurement cycles of direct institutional sales, while creating internal champions who can drive school-level licensing decisions from within

10.3 KEY TAKEAWAYS

- AI output validation is non-negotiable for production applications. Without Pydantic enforcement, ExamCraft would be unreliable. With it, reliability exceeds 97%.
- Workflow design matters more than AI capability. The structured wizard UI and section-level configuration produce consistently usable output where raw ChatGPT prompting consistently fails educators.
- Pedagogical authenticity accelerates adoption. Bloom's Taxonomy integration makes ExamCraft immediately credible to professional educators who are skeptical of AI tools.
- Document security is a feature, not a luxury. PDF encryption and PIN systems are not typical features of developer-side AI tools; for institutional educators, they are prerequisites.
- Infrastructure costs are not a barrier. The total operational cost of under $65/month for early production deployment makes this viable as a student startup.

10.4 FINAL REMARKS

ExamCraft.ai represents a complete, functional demonstration of AI-native software engineering for the EdTech domain, delivering measurable value to educators while establishing a technically sound foundation for continued product development and commercial growth. The limitations identified — particularly the absence of a manual editing UI and the context window constraint for large documents — are well-understood engineering problems with clear solution paths in the roadmap.

The project team has produced not merely an academic exercise but a genuinely deployable product with real commercial potential. With the completion of the v1.1 Rich Text Editor and RAG integration milestones, ExamCraft.ai will be positioned for a focused go-to-market launch targeting the Indian EdTech market's 9+ million active school teachers.

CHAPTER 11

APPENDICES

- **BUSINESS MODEL**
- **PRODUCT DEPLOYMENT CONFIGURATION**
- **API REFERENCE**
- **ENVIRONMENT VARIABLES REFERENCE**

11.1 BUSINESS MODEL

ExamCraft.ai operates on a B2B2C SaaS business model targeting educational institutions and individual educators. The model is structured around three revenue tiers:

Tier	Price	Features	Target User
Free / Trial	$0/month	3 generations/month, watermarked PDFs, no PDF encryption, standard question types	New users, evaluation
Educator Pro	$12–15/month	Unlimited generations, DOCX export, PDF encryption, custom templates, answer keys, full PIN security	Individual teachers & tutors
Institutional	$150+/month	Bulk teacher seats, branded PDFs with school logo, shared question bank, centralized analytics, priority support	Schools, coaching institutes

Revenue Projections

Break-even requires approximately 20 Pro subscribers covering estimated monthly infrastructure costs of $250. At 500 Pro subscribers, monthly recurring revenue (MRR) reaches approximately $7,500. Securing 20 institutional accounts at $150/month yields an additional $3,000 MRR with significantly lower churn rates than individual subscriptions.

Go-To-Market Strategy

Phase 1 targets individual tech-savvy teachers through social media content marketing (LinkedIn, Twitter/X) and teaching community platforms. A 'Hero Teacher' program provides lifetime Pro access in exchange for institutional referral and public demonstration. Phase 2 targets coaching institutes (JEE/NEET prep centers) requiring high-volume daily paper generation. Phase 3 pursues school and university district licensing with longer sales cycles but higher contract values and lower churn.

11.2 PRODUCT DEPLOYMENT CONFIGURATION

11.2.1 Backend Deployment

Base Image	Python 3.12 slim
WSGI Server	Waitress (production-grade, cross-platform Python WSGI server)

Entry Point	wsgi.py imports and serves the Flask app object
Port	5000 (internal), mapped to host port 5000
Environment	Variables loaded from .env.local via python-dotenv
Required Env Vars	GROQ_API_KEY, SUPABASE_URL, SUPABASE_SERVICE_ROLE_KEY, ALLOWED_ORIGINS
Uploads Volume	/uploads directory mounted as Docker volume for persistent file storage
Rate Limits	Flask-Limiter: 10 req/min for /upload, 5 req/min for /generate
Logging	RotatingFileHandler to server.log (5MB max, 3 backup files)
MAX_CONTENT_LENGTH	50MB (Flask config: 50 * 1024 * 1024)

11.2.2 Frontend Deployment

Build Tool	Vite — produces optimized static bundle in dist/
Production Serving	Static bundle deployable to Netlify, Vercel, or Nginx
API Communication	Backend URL configured in Frontend/.env.local as VITE_API_URL
Authentication	Supabase JS client handles JWT session management client-side
Required Env Vars	VITE_SUPABASE_URL, VITE_SUPABASE_ANON_KEY, VITE_API_URL

11.3 API REFENRENCE

The ExamCraft.ai backend exposes the following RESTful API endpoints. All endpoints except the home route require a valid Supabase JWT Bearer token in the Authorization header (Authorization: Bearer <token>).

Method & Endpoint	Rate Limit	Auth	Description
GET /	None	None	Health check. Returns API status, message, and version string.
POST /upload	10/min	Required	Upload study material PDF. Returns file_id, db_id, and content_preview.

POST /generate	5/min	Required	Generate exam paper from content/file_ids and configure_data. Returns full ExamPaper JSON with paper_id.
POST /download	None	Required	Generate and download exam in PDF or DOCX format. Accepts format ('pdf'|'docx') and optional password.
POST /regenerate-question	None	Required	Regenerate a single question. Accepts question_data, paper_id, source_content.
GET /api/profile	None	Required	Retrieve authenticated user's profile record.
PUT /api/profile	None	Required	Update profile fields (full_name, institution_name, role, work_email, exam_pin_hash).
POST /api/verify-exam-password	None	Required	Verify 6-digit exam PIN against stored hash. Returns {valid: true|false}.
GET /api/exam_papers	None	Required	List all exam papers for authenticated user with metadata.
GET /api/exam_papers/<id>	None	Required	Retrieve single exam paper including full content JSON.
GET /api/materials	None	Required	List all uploaded materials for authenticated user.
DELETE /api/materials/<id>	None	Required	Delete material record and associated Supabase Storage file.

11.4 ENVIRONMENT VARIABLES REFERENCE

Variable	Location	Purpose
GROQ_API_KEY	Backend .env.local	Groq AI API authentication key
SUPABASE_URL	Backend .env.local	Supabase project endpoint URL
SUPABASE_SERVICE_ROLE_KEY	Backend .env.local	Server-side Supabase auth bypass key
ALLOWED_ORIGINS	Backend .env.local	Comma-separated CORS whitelisted frontend origins
FLASK_DEBUG	Backend .env.local	Enable/disable Flask debug mode (true|false)
VITE_SUPABASE_URL	Frontend .env.local	Supabase URL for client-side SDK
VITE_SUPABASE_ANON_KEY	Frontend .env.local	Supabase anonymous key for client SDK
VITE_API_URL	Frontend .env.local	Backend Flask API base URL

BIBLIOGRAPHY

The following references informed the design, technology choices, and theoretical framework of the ExamCraft.ai platform:

- Bloom, B.S., Engelhart, M.D., Furst, E.J., Hill, W.H., & Krathwohl, D.R. (1956). Taxonomy of educational objectives: The classification of educational goals. Handbook I: Cognitive domain. David McKay Company.
- Anderson, L.W., & Krathwohl, D.R. (Eds.). (2001). A taxonomy for learning, teaching, and assessing: A revision of Bloom's taxonomy of educational objectives. Longman.
- Du, X., Shao, J., & Cardie, C. (2017). Learning to ask: Neural question generation for reading comprehension. Proceedings of ACL 2017.
- Heilman, M., & Smith, N.A. (2010). Good question! Statistical ranking for question generation. Proceedings of ACL 2010.
- Vaswani, A., et al. (2017). Attention is all you need. Advances in Neural Information Processing Systems (NeurIPS 2017).
- Brown, T., et al. (2020). Language models are few-shot learners. Advances in Neural Information Processing Systems 33 (NeurIPS 2020).
- Flask Documentation. (2024). Flask 3.0 — A Micro Web Framework for Python. https://flask.palletsprojects.com/
- React Documentation. (2024). React 18 Official Documentation. https://react.dev/
- Supabase Documentation. (2024). Supabase — The Open Source Firebase Alternative. https://supabase.com/docs
- Groq Documentation. (2024). Groq API Reference and Model Specifications. https://console.groq.com/docs/
- Pydantic Documentation. (2024). Pydantic V2 — Data Validation using Python Type Hints. https://docs.pydantic.dev/
- Van Rossum, G., & Python Software Foundation. (2024). Python 3.12 Documentation. https://docs.python.org/3.12/
- Tailwind CSS Documentation. (2024). Tailwind CSS v3 — Utility-First CSS Framework. https://tailwindcss.com/docs
- ReportLab. (2024). ReportLab PDF Generation User Guide. https://www.reportlab.com/docs/reportlab-userguide.pdf
- Python-docx Documentation. (2024). python-docx — Create and Update Microsoft Word Files. https://python-docx.readthedocs.io/
- Meta AI. (2024). Llama 3.3 Technical Report. https://ai.meta.com/research/publications/